

1 ALMA MATER STUDIORUM · UNIVERSITY OF BOLOGNA
2

3
4 School of Science
5 Department of Physics and Astronomy
6 Master Degree in Physics

7

Dataset Generation for the Training of 8 Neural Networks Oriented toward 9 Histological Image Segmentation

10 **Supervisor:**
Dr. Enrico Giampieri

Co-supervisor:
Dr. Nico Curti

11
Submitted by:
Alessandro d'Agostino

12 Academic Year 2019/2020

¹⁴ Abstract

¹⁵ The project is inspired by an actual problem of timing and accessibility in the analysis of
¹⁶ histological samples in the health-care system. In this project, I face the problem of syn-
¹⁷ thetic histological image generation for the purpose of training Neural Networks for the
¹⁸ segmentation of real histological images. The collection of real histological human-labeled
¹⁹ samples is a very time consuming and expensive process and often is not representative
²⁰ of healthy samples, for the intrinsic nature of medical analysis. The method I propose is
²¹ based on the replication of the traditional specimen preparation technique in a virtual
²² environment. The first step is the creation of a 3D virtual model of a region of the target
²³ human tissue. The model should encapture all the key features of the tissue, and the
²⁴ richer it is the better will be the yielded result. The second step is to perform a sampling
²⁵ of the model through a virtual tomography process, which produces a first completely
²⁶ labeled image of the section. This image is then processed with different tools to achieve
²⁷ a histological-like aspect. The most significant aesthetical adjustment is given by the
²⁸ action of a style transfer neural network that implants the typical histological visual tex-
²⁹ ture on the synthetic image. This procedure is presented in detail for two specific models
³⁰ of human tissue: one of pancreatic tissue and one of dermal tissue. The two resulting
³¹ images compose a pair of images perfectly suitable for a supervised learning technique.
³² The generation process is completely automatized and does not require the intervention
³³ of any human operator, hence it can be used to produce arbitrary large datasets. The
³⁴ synthetic images are inevitably less complex than the real samples and they offer an
³⁵ easier segmentation task to solve for the NN. However, the synthetic images are very
³⁶ abundant, and the training of a NN can take advantage of this feature, following the
³⁷ so-called curriculum learning strategy.

³⁸ Table of Contents

³⁹	Introduction	6
⁴⁰	1 Theoretical Background	12
⁴¹	1.1 Tissues Microanatomy and Tumor Evidences	12
⁴²	1.1.1 Pancreas Microanatomy	12
⁴³	1.1.2 Skin Microanatomy	13
⁴⁴	1.2 Histological Images Digitalization	15
⁴⁵	1.2.1 Slides Preparation for Optic Microscopic Observation	17
⁴⁶	1.3 Introduction to Deep Learning	19
⁴⁷	1.3.1 Perceptrons and Multilayer Feedforward Architecture	19
⁴⁸	1.3.2 Training Strategies in Deep Learning	21
⁴⁹	1.3.3 Training Algorithms - Error Back-Propagation	23
⁵⁰	1.4 Deep Learning-Based Segmentation Algorithms	27
⁵¹	1.4.1 State of the Art on Deep Learning Segmentation	28
⁵²	1.4.2 Image Segmentation Datasets	33
⁵³	2 Technical Tools for Model Development	36
⁵⁴	2.1 Quaternions	36
⁵⁵	2.2 Parametric L-Systems	38
⁵⁶	2.3 Voronoi Tassellation	40
⁵⁷	2.4 Saltelli Algorithm - Randon Number Generation	42
⁵⁸	2.5 Planar Section of a Polyhedron	44
⁵⁹	2.6 Perlin Noise	46
⁶⁰	2.7 Style-Transfer Neural Network	47
⁶¹	2.8 Working Environment	50
⁶²	3 Tissue Models Development	52

63	3.1	Pancreatic Tissue Model	52
64	3.1.1	2D Ramification	52
65	3.1.2	Expansion to 3D	54
66	3.1.3	Subdivision in Cells	55
67	3.1.4	Cells Identity Assignment	58
68	3.2	Dermal Tissue Model	59
69	3.3	Synthetic Images Production	62
70	3.3.1	Sectioning Process	62
71	3.3.2	Appearance Makeover	64
72	Conclusions		71
73	Bibliography		74

⁷⁵ Introduction

⁷⁶ In the last decades, the development of Machine Learning (ML) and Deep Learning (DL)
⁷⁷ techniques has contaminated every aspect of the scientific world, with interesting results
⁷⁸ in many different research fields. The biomedical field is no exception to this and a
⁷⁹ lot of promising applications are taking form, especially as Computer-Aided Detection
⁸⁰ (CAD) systems which are tools for the support for physicians during the diagnostic
⁸¹ process. Medical doctors and the healthcare system in general collect a huge amount of
⁸² data from patients during all the treatment, screening, and analysis activities in many
⁸³ different shapes, from anographical data to blood analysis to clinical images.

⁸⁴ In fact in medicine, the study of images is ubiquitous and countless diagnostic proce-
⁸⁵ dures rely on it, such as X-ray imaging (CAT), nuclear imaging (SPECT, PET), Magnetic
⁸⁶ resonance, and visual inspection of histological specimens after biopsies. The branch of
⁸⁷ artificial intelligence in the biomedical field that handles image analysis to assist physi-
⁸⁸ cians in their clinical decisions goes under the name of Digital Pathology Image Analysis
⁸⁹ (DPIA). In this thesis work, I want to focus on some of the beneficial aspects intro-
⁹⁰ duced by DPIA in the histological images analysis and some particular issues in the
⁹¹ development of DL models able to handle this kind of procedure.

⁹² Nowadays the great majority of analysis of histological specimens occurs through
⁹³ visual inspection, carried out by highly qualified experts. Some analysis, as cancer
⁹⁴ detection, requires the ability to distinguish if a region of tissue is healthy or not with high
⁹⁵ precision in very wide specimens. This kind of procedure is typically very complex and
⁹⁶ requires prolonged times of analysis besides substantial economic efforts. Furthermore,
⁹⁷ the designated personnel for this type of analysis is often limited, leading to delicate issues
⁹⁸ of priority assignment while scheduling analysis, based on the estimated patient's clinical
⁹⁹ development. Some sort of support to this analysis procedure is therefore necessary.

¹⁰⁰ The problem of recognizing regions with different features within an image and de-
¹⁰¹ tect their borders is known in computer vision as the segmentation task, and it's quite
¹⁰² widespread with countless different applications, allowing a sort of automatic image in-
¹⁰³ terpretation. In ML the segmentation problem is usually faced as a supervised task,
¹⁰⁴ hence the algorithm in order to be trained properly requires an appropriate quantity
¹⁰⁵ of pre-labeled images, from which learn the rules through which distinguish different
¹⁰⁶ regions. This means that the development of segmentation algorithms for a specific ap-

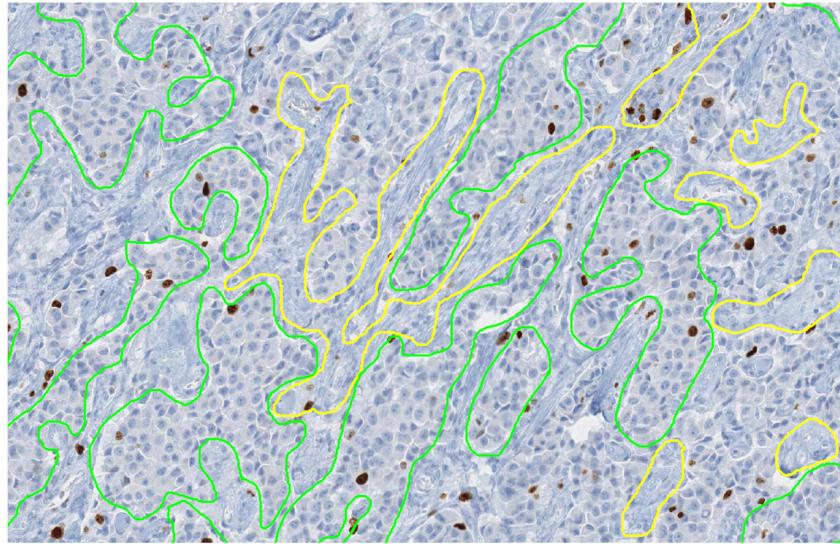


Figure 1: Interleaving of tumor (green annotation) and non-tumor (yellow annotation) regions [28].

107 plication, as would be the one on histological images, would require a lot of starting
108 material, previously analyzed from the same qualified expert encharged of the visual
109 inspection mentioned before. A human operator thus is required to manually track the
110 boundaries, for example, between healthy and tumoral regions within a sample of tissue
111 and to label them with their identity, as in Figure 1. The more the algorithm to train
112 is complex the more starting material is required to adjust the model’s parameters and
113 reach the desired efficacy.

114 The latest developed segmentation algorithms are based on DL techniques, hence
115 based on the implementation of intricated Neural Networks (NN) which process the
116 input images and produce the corresponding segmentation. Those models are typically
117 very complex, with millions of parameters to adjust and tune, therefore they need a
118 huge amount of pre-labeled images to learn their segmentation rules. This need for data
119 is exactly the main focus of my thesis work. The shortage of ground truth images is
120 indeed one of the toughest hurdles to overcome during the development of DL-based
121 algorithms. Another important aspect to bear in mind is the quality of the ground truth
122 material. It’s impossible for humans to label boundaries of different regions with pixel-
123 perfect precision, while for machines the more precise is the input the more effective is
124 the resulting algorithm.

125 Different approaches have already been explored to overcome this problem, and they
126 are mainly based on the generation of synthetic labeled data to use during the training
127 phase. Some techniques achieve data augmentation manipulating already available
128 images and then generating *new* images, but as we will see this approach suffers from

129 different issues. Here, I want to make an overview of some other interesting works on
130 the generation of synthetic histological images, which have followed completely different
131 paths and strategies from mine.

132 The first work I want to cite is a work from Ben Cheikh *et al.* from 2017 [9]. In
133 this work, they present a methodology for the generation of synthetic images of different
134 types of breast carcinomas. They propose a method completely based on two-dimensional
135 morphology operation, as successive image dilations and erosions. With the modulation
136 of a very restricted number of parameters, regulating the abundance of the objects, their
137 distribution in the image, and their shapes they are able to reconstruct realistic images
138 reflecting different histological situations. In Figure 2 is shown an example of generated
139 material besides a real histological H&E stained sample. Starting from a generated
140 segmentation mask which defines the *tumoral pattern* the production of the synthetic
141 image passes through successive steps, as the generation of characteristic collagen fibers
142 around the structure, the injection of all the immune system cells, and some general final
143 refinements.

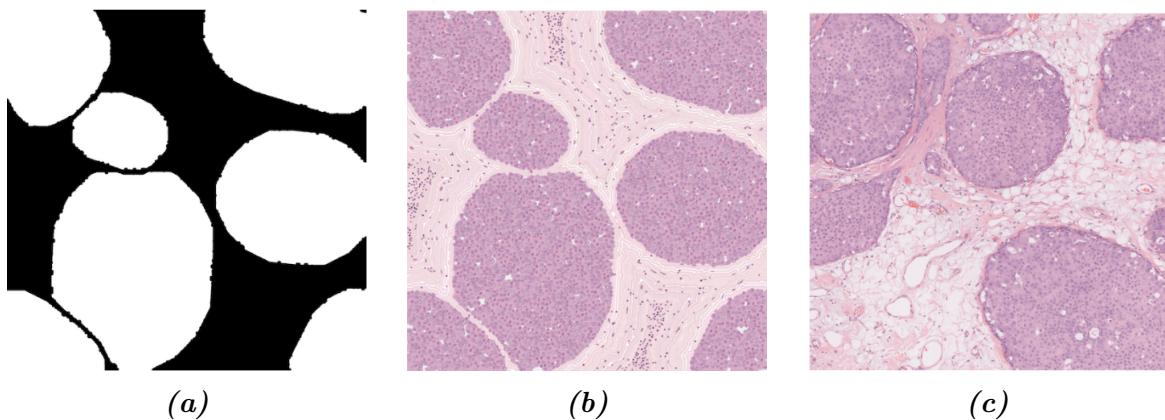


Figure 2: Example of generated tumoral pattern (left), which acts as segmentation mask, of generated image (center) and a real example of the tissue to recreate, from [9].

144 The second work I want to mention is based on a DL-base technique, which ap-
145 proaches synthetic image generation using a specific cGAN architecture inspired to the
146 “U-net” [38] model, as will be described in Figure 1.15 in section 1.4.1. This model
147 works with Ki67 stained samples of breast cancer tissue, and it is able to generate high-
148 fidelity images starting from a given segmentation mask. Those starting segmentation
149 masks tough are obtained through the processing of other real histological samples, via
150 a nuclei-detection algorithm. The differences between real and synthetic samples are
151 imperceptible, and the material generated in this work has effectively fooled experts,
152 who qualified it has indistinguishable from the real one. In Figure 3 an example of a real
153 image, a generated one, and their corresponding segmentation mask.

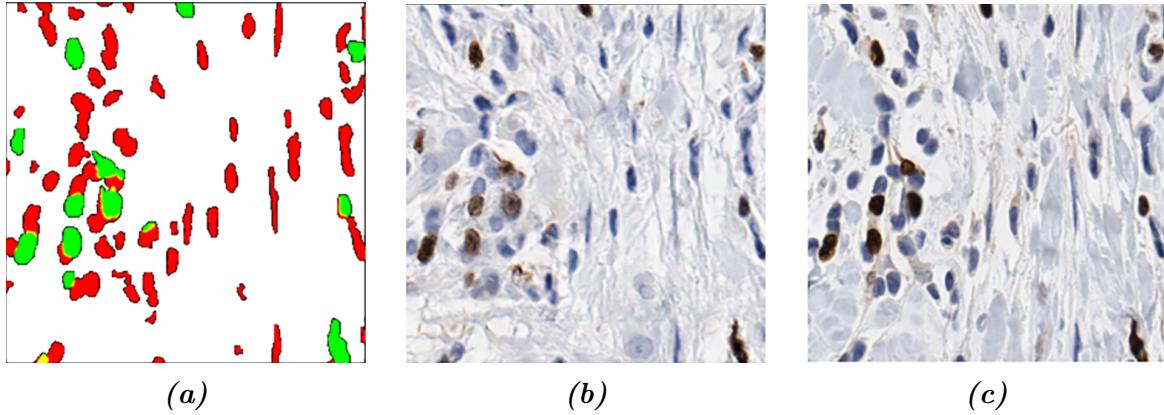


Figure 3: Example of generated tumoral pattern (left), which acts as segmentation mask, of generated image (center) and a real example of the tissue to recreate, from [38].

Both of the two before-mentioned strategies produces realistic (or even perfect) results, but they are based onto considerations and analysis limited only to the aspect of the images. In the first work, the segmentation mask is produced in an almost full-random way, while in the second the segmentation mask is extracted starting from an actual real histological samples. The target of the present work instead lies in between those two approaches, and wants to produce randomized new images following a plausible modelization based on physical and histological considerations.

The technique I propose in this work follows a generation from scratch of entire datasets suitable for the training of new algorithms, based on the 3D modelization of a region of human tissue at the cellular level. The entire traditional sectioning process, which is made on real histological samples, is recreated virtually on this virtual model. This yields pairs of synthetic images with their corresponding ground-truth. Using this technique one would be able to collect sufficient material for the training (the entire phase or the preliminary part) of a model, overcoming the shortage of hand-labeled data. The 3D modeling of a region of particular human tissue is a very complex task, and it is almost impossible to capture all the physiological richness of a histological system. The models I implemented thus are inevitably less sophisticated respect the target biological structures. I'll show two models: one of pancreatic tissue and another of dermal tissue, besides all the tools I used and the choices I made during the designing phase.

Furthermore, since the image production passes through a wide and elaborated model, the resulting images contains a new level of semantic information that would not be encapturable otherwise. From the modelization is possible to reflect on the segmentation mask image the relationship information between nuclei and their belonging cells, or the basins of blood irrigation corresponding to every blood vessels and many other informations about the interrelationship of depicted elements. Moreover, achieving the right

mastery of the modelization it is possible to reproduce different physiological states of the tissue like the healthy *standard* configuration or different pathological situations, which reflect themselves as particular arrangements at cellular level. This aspect is of great interest, in fact there is a strong lack of real histological samples of healthy tissues samples, given the intrinsic nature of medical analysis. Biopsies are in general invasive procedure, and are typically performed only when there is a concern for a pathological situation. Unless an erroneous evaluation, they typically collect a sample of non-healthy tissue. This method thus allows to collect an arbitrary number of samples in every interesting modeled condition, overcoming the problem of under-represented conditions.

In order to present logically all the steps of my work the thesis is organized in chapters as follows:

1. Theoretical Background .

In this chapter, I will describe how real histological images are obtained and their digitalization process works. Afterward, I will introduce the reader to the Deep Learning framework, explaining the key elements of this discipline and how they work. Finally, I will dedicate a section to the image segmentation problem, and the state of the art of segmentation DL-based algorithms, with particular attention to the applications in the bio-medical field.

2. Technical Tools for Model Development .

I will dedicate this chapter to the thorough description of every technical tool I needed during the designing phase of this project. The development has required the harmonization of many different technologies and mathematical tools, some of which not so popular like quaternions, quasi-random number generation, Voronoi decomposition, and style transfer neural networks. In this chapter I will describe also the specialized algorithm I devised and implemented for the sectioning of an arbitrary polyhedron, which is the key element for the correct working of the virtual tomography technique described by this thesis work. As a conclusion for this chapter, I will describe the working environment I built for developing this project and I will mention all the code libraries I employed in my work.

3. Tissue Models Development .

This third chapter is the heart of the project. I will describe in detail all the steps necessary to create the two models, one of pancreatic tissue and the other of dermal tissue, and how I am able to produce the synthetic images. The first section is devoted to the modeling of the histological structures, while the second is entirely dedicated to the sectioning process and the subsequent refinements to the images.

²¹⁶ **Chapter 1**

²¹⁷ **Theoretical Background**

²¹⁸ In this first chapter, I will depict the theoretical context of the work. Section 1.1 con-
²¹⁹ tains a brief introduction to the anatomy of the two particular tissues under the attention
²²⁰ of this project. Section 1.2 will be dedicated to histological images, and the different
²²¹ techniques used to prepare the samples to analyze. Histological images recover a funda-
²²² mental role in medicine and are the pillar of many diagnosis techniques. This discipline
²²³ borns traditionally from the optical inspection of the tissue slides using a microscope,
²²⁴ and it is gradually developing and improving with the advent of computers and digital
²²⁵ image processing. It is important tough to understand how the samples are physically
²²⁶ prepared, the final target of this work is in fact the virtual reconstruction of this process.
²²⁷ In section 1.3 I will introduce the Deep Learning framework and describe how a Neural
²²⁸ Network works and actually learns. The most advanced techniques for the automatic
²²⁹ image processing implement Deep Learning algorithm, and understanding the general
²³⁰ rules behind this discipline is crucial for a good comprehension of this work. In section
²³¹ 1.4 I will discuss in particular the problem of image segmentation and how it is tackled
²³² with different Neural Network architectures, showing what it is the state of the art of
²³³ this research field.

²³⁴ **1.1 Tissues Microanatomy and Tumor Evidences**

²³⁵ **1.1.1 Pancreas Microanatomy**

²³⁶ The Pancreas is an internal organ of the human body, part of both the digestive system
²³⁷ and the endocrine system. It acts as a gland with both endocrine and exocrine functions,
²³⁸ and it is located in the abdomen behind the stomach. Its main endocrine duty is the
²³⁹ secretion of hormones like insulin and glucagon which are responsible for the regulation
²⁴⁰ of sugar levels in the blood. As a part of the digestive system instead it acts as an exocrine
²⁴¹ gland secreting pancreatic juice. The majority of pancreatic tissue has a digestive role,

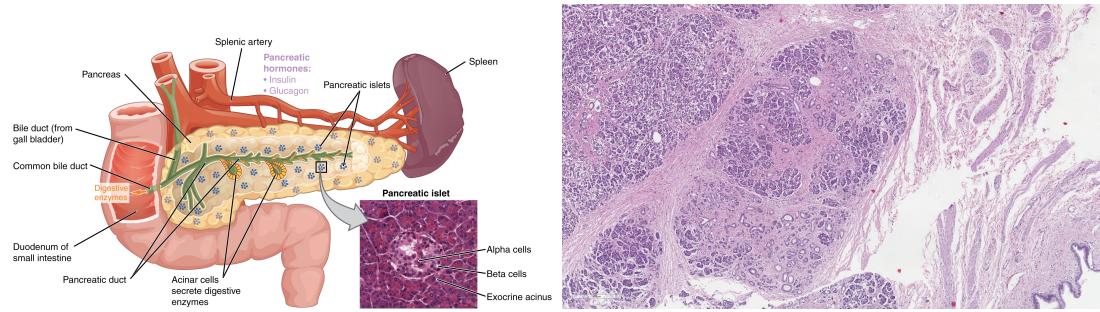


Figure 1.1: (left) A picture of pancreas' structure in its physiological context. In this picture is clearly visible the macroscopic structure and the glandular organization at microscopic level, and how it reflects in the histological sample. (right) A real pancreatic tissue sample with H&E staining.

and the cells with this role form clusters (*acini*) around the small pancreatic ducts. The acinus secrete inactive digestive enzymes called zymogens into the small intercalated ducts which they surround, and then in the pancreatic blood vessels system [26]. In Figure 1.1 is shown a picture of the pancreas, with its structure and its placement in the human body.

All the tissue is actually rich in other important elements as the islets of Langerhans, and sporadic connective tissue all over the structure, which are clearly visible in the traditional histological specimens.

1.1.2 Skin Microanatomy

Skin is the layer of soft, flexible outer tissue covering the body of a vertebrate animal, with the three main functions of protection, regulation, and sensation. Mammalian skin is composed of two primary layers: the epidermis, which provides waterproofing and serves as a barrier to infection, and the dermis, which serves as a location for the appendages of skin. The epidermis is composed of the outermost layers of the skin. It forms a protective barrier over the body's surface, responsible for keeping water in the body and preventing pathogens from entering, and is a stratified squamous epithelium, composed of proliferating basal and differentiated suprabasal keratinocytes. The dermis is the layer of skin beneath the epidermis that consists of connective tissue and cushions the body from stress and strain. The dermis provides tensile strength and elasticity to the skin through an extracellular matrix composed of collagen fibrils, microfibrils, and elastic fibers, embedded in hyaluronan and proteoglycans.

20c2734

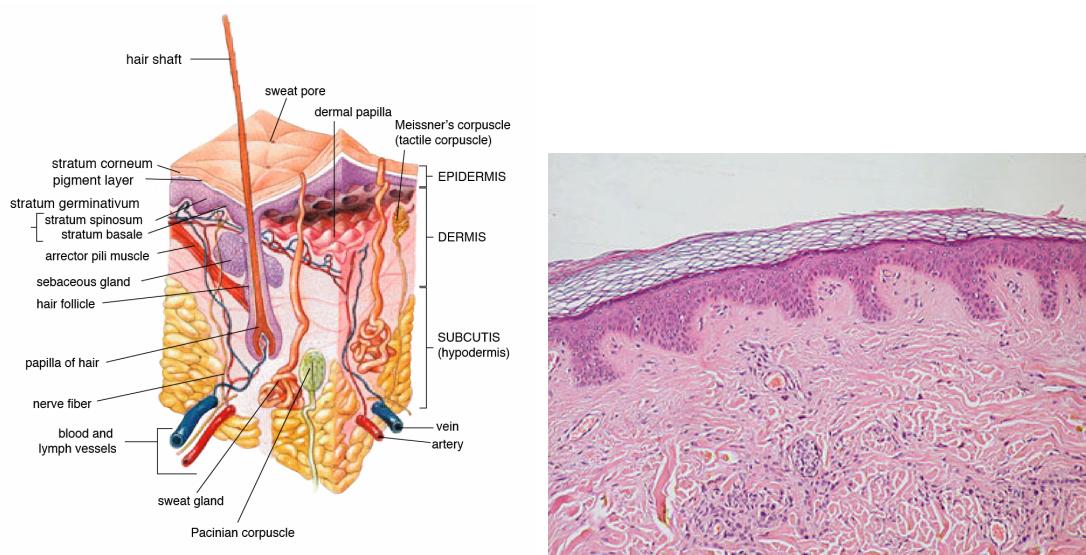


Figure 1.2: (left) Microanatomical description of a region of dermal tissue and all the interesting elements present in cutis, and subcutaneous layer. (right) An actual histological specimen from a sample of dermal tissue.

264 1.2 Histological Images Digitalization

265 Modern histopathology is essentially based on the careful interpretation of microscopic
266 images, with the intention of correctly diagnose patients and to guide therapeutic de-
267 cisions. In the last years, thanks to the quick development of scanning techniques and
268 image processing, the discipline of histology have seen radical improvements: the main
269 of which undoubtedly is the passage from the microscope's oculars to the computer's
270 screen. This digitalization process has brought several advantages, that were previously
271 impossible in classical histology, like telepathology and remote assistance in diagnosis
272 processes, the integration with other digitalized clinical workflows, and patients' history,
273 and most importantly the opening to applications of artificial intelligence. The name
274 Whole Slide Imaging (WSI) refers to the modern virtual microscopy discipline, which
275 consists of scanning a complete microscope slide and creating a single high-resolution
276 digital file. This is commonly achieved by capturing many small high-resolution image
277 tiles or strips and then montaging them to create a full image of a histological section.
278 The four key steps of this process are image acquisition (scansion), editing, and on-screen
279 image visualization.

280 In the field of Digital Pathology (DP) an essential concept in image understanding
281 is the magnification factor, which indicates the scale of representation of the image and
282 allows dimension referencing. This factor is usually indicated as the magnification power
283 of the microscope's lenses used during the analysis. After the digitalization process,
284 this original magnification factor is prone to change, depending on the resolution of the
285 visualization screen. Therefore, image resolution is measured in μm per pixel, and it is set
286 by the different composition of the acquisition chain, as the optical sensor and the lenses.
287 Histological scanner are usually equipped with $20\times$ or $40\times$ objectives, which correspond
288 to 0.5 and $0.25 mm/pixel$ resolution values. Lenses with $20\times$ magnification factor are
289 the most suitable for the great majority of histopathological evaluations, and it is the
290 golden standard for scansions, for its good trade-off between image quality and time of
291 acquisition. Scansions with $40\times$ magnification could increase four-fold acquisition and
292 processing time, final file's dimension, and storage cost. A single WSI image, acquired
293 with $20\times$ will occupy more than 600 MB alone.

294 Despite the WSI is a relatively mature discipline, it still struggles to integrate itself in
295 the standard primitive diagnosis phase in histopathological laboratories. This is primar-
296 ily due to some disadvantages, like images'resolution, image compression's artifacts, and
297 auto-focusing algorithms, which plays a key role in the specimen interpretation. Fur-
298 thermore, the scansion of histological samples is an additional step in the analysis which
299 takes time. Despite the technological improvements the average time for the acquisition
300 of a sample is around 5/10 minutes, depending on the number of slices in the slide, for
301 just a single level of magnification. While in traditional histology, the pathologist has
302 access to all the magnification levels at the same time. The real advantage, in fact, is in
303 the long term. Once the images have been acquired they can be archived and consulted

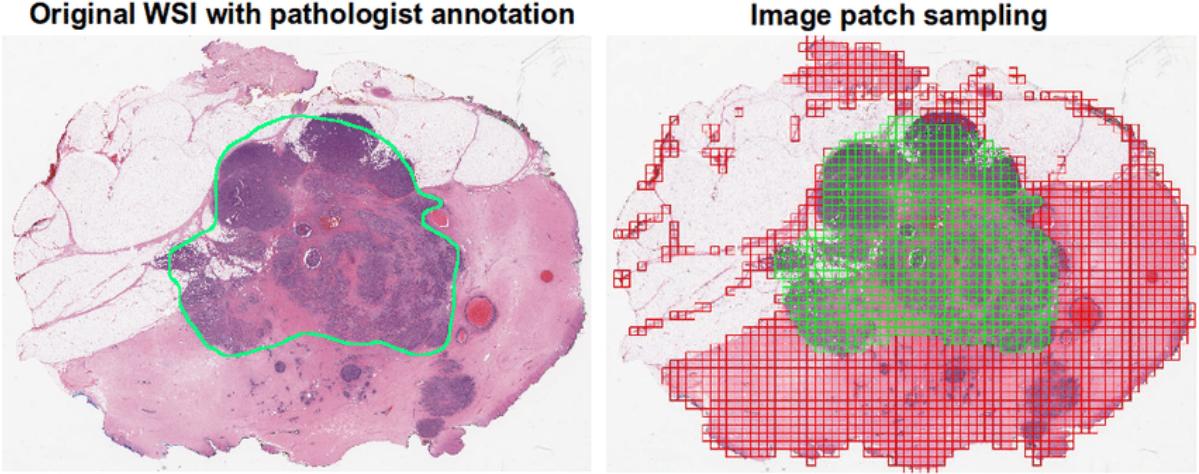


Figure 1.3: An example of whole slide image, with its grid decomposition in patches. It is visible the correspondence between a region of interest manually annotated and the patches that matches that region. From [13]

304 remotely almost instantaneously, helping clinical analysis and allowing remote assistance
 305 (telemedicine). Furthermore, the images now can be processed by artificial intelligence
 306 algorithms, allowing the application of technologies like Deep Learning which could rev-
 307 olutionize the research field, as already has been on many different disciplines in the
 308 scientific world.

309 In order to allow to automatically process, such big images as the ones obtained
 310 through WSI, it is necessary to subdivide them in smaller patches. The dimension of
 311 which should be big enough to allow interpretation and to preserve a certain degree
 312 of representability of the original image. In Figure 1.3 is shown an example of whole
 313 slide image, with its grid decomposition in patches. If the patches are too small, it
 314 should be over-specified for a particular region of tissue, loosing its general features.
 315 This could lead the learning algorithm to misinterpretation. However, this is not an
 316 exclusive limit of digital pathology, for a human pathologist would be impossible too to
 317 make solid decisions on a too limited sample of tissue. After the subdivision in patches, a
 318 typical process for biomedical images is the so-called *data augmentation* of images, that
 319 is the process of creating re-newed images from the starting material through simple
 320 geometrical transformations, like translation, rotation, reflection, zoom in/out.

321 The analysis of histological images usually consists in detecting the different com-
 322 ponents in the samples and to recognise their arrangement as an healthy or pathological
 323 pattern. It is necessary to recognize every sign of vitality of the cells, evaluating the
 324 state of the nucleus. There are many additional indicators to consider like the presence
 325 of inflammatory cells or tumoral cells. Furthermore, samples taken from different part
 326 of the human body present completely different characteristics, and this increase greatly

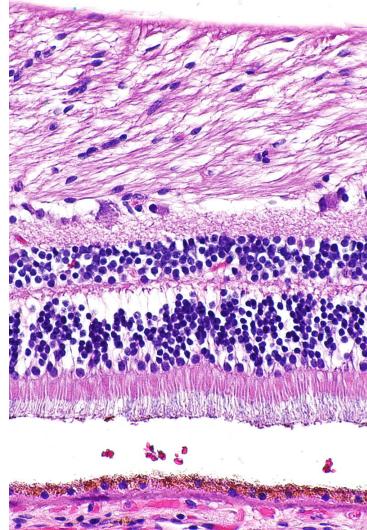


Figure 1.4: A sample of tissue from a retina (a part of the eye) stained with hematoxylin and eosin, cell nuclei stained blue-purple and extracellular material stained pink.

327 the complexity of the analysis. A reliable examination of a sample thus require a care-
328 ful inspection made by an high qualified expert. The automatization of this procedure
329 would be extremily helpful, giving an increadible boost both in timing and in accessibil-
330 ity. However, this is not a simple task and in section 1.4.1 I will show some actual model
331 for biomedical image processing in detail.

332 1.2.1 Slides Preparation for Optic Microscopic Observation

333 In modern, as in traditional, histology regardless on the final support of the image the
334 slide has to be physically prepared, starting from the sample of tissue. The sample and
335 slide preparation is a crucial step for histological or cytological observation. It is essential
336 to highlight what needs to be observed and to *immobilize* the sample at a particular point
337 in time and with characteristics close to those of its living state. There are five key steps
338 for the preparation of samples [4]:

- 339 1) **Fixation** is carried out immediately after the removal of the sample to be observed.
340 It is used to immobilize and preserve the sample permanently in as life-like state
341 as possible. It can be performed immersing the biological material in a formalin
342 solution or by freezing, so immersing the sample in a tissue freezing medium which
343 is then cooled in liquid nitrogen.
- 344 2) **Embedding** if the sample has been stabilized in a fixative solution, this is the sub-
345 sequent step. It consists in hardening the sample in a paraffin embedding medium,

346 in order to be able to carry out the sectioning. It is necessary to dehydrate the
347 sample beforehand, by replacing the water molecules in the sample with ethanol.

348 **3) Sectioning** Sectioning is performed using microtomy or cryotomy. Sectioning is an
349 important step for the preparation of slides as it ensures a proper observation of the
350 sample by microscopy. Paraffin-embedded samples are cut by cross section, using a
351 microtome, into thin slices of $5 \mu m$. Frozen samples are cut using a cryostat. The
352 frozen sections are then placed on a glass slide for storage at $-80^{\circ}C$. The choice of
353 these preparation conditions is crucial in order to minimize the artifacts. Paraffin
354 embedding is favored for preserving tissues; freezing is more suitable for preserving
355 DNA and RNA and for the labeling of water-soluble elements or of those sensitive
356 to the fixation medium.

357 **4) Staining** Staining increases contrasts in order to recognize and differentiate the dif-
358 ferent components of the biological material. The sample is first deparaffinized and
359 rehydrated so that polar dyes can impregnate the tissues. The different dyes can
360 thus interact with the components to be stained according to their affinities. Once
361 staining is completed, the slide is rinsed and dehydrated for the mounting step.

362 Hematoxylin and eosin stain (H&E) is one of the principal tissue stains used in
363 histology [44], and it is the most widely used stain in medical diagnosis and is often
364 the gold standard [35]. H&E is the combination of two histological stains: hematoxylin
365 and eosin. The hematoxylin stains cell nuclei blue, and eosin stains the extracellular
366 matrix and cytoplasm pink, with other structures taking on different shades, hues, and
367 combinations of these colors. An example of H&E stained is shown in Figure 1.4, in
368 which we can see the typical colour palette of an histological specimen.

369 1.3 Introduction to Deep Learning

370 Deep Learning is part of the broader framework of Machine Learning and Artificial
371 Intelligence. Indeed all the problems typically faced using ML can also be addressed
372 with DL techniques, for instance, regression, classification, clustering, and segmentation
373 problems. We can think of DL as a universal methodology for iterative function ap-
374 proximation with a great level of complexity. In the last decades, this technology has
375 seen a frenetic diffusion and an incredible development, thanks to the always increasing
376 available computational power, and it has become a staple tool in all sorts of scientific
377 applications.

378 1.3.1 Perceptrons and Multilayer Feedforward Architecture

379 Like other artificial learning techniques, DL models aim to *learn* a relationship between
380 some sort of input and a specific kind of output. In other words, approximating nu-
381 merically the function that processes the input data and produces the desired response.
382 For example, one could be interested in clustering data in a multidimensional features
383 space, or the detection of objects in a picture, or text manipulation/generation. The
384 function is approximated employing a greatly complex network of simple linear and non-
385 linear mathematical operations arranged in a so-called Neural Network (typically with
386 millions of parameters). The seed idea behind this discipline is to recreate the function-
387 ing of actual neurons in the human brain: their entangled connection system and their
388 “ON/OFF” behavior [41].

389 The fundamental unit of a neural network is called perceptron, and it acts as a digital
390 counterpart of a human neuron. As shown in Figure 1.5 a perceptron collects in input a
391 series on n numerical signals $\vec{x} = 1, x_1, \dots, x_n$ and computes a linear weighted combination
392 with the weights vectors $\vec{w} = w_0, w_1, \dots, w_n$, where w_0 is a bias factor:

$$f(\vec{x}, \vec{w}) = \chi(\vec{x} \cdot \vec{w}). \quad (1.1)$$

393 The results of this linear combination are given as input to a non-linear function $\chi(x)$
394 called the activation function. Typical choices as activation function are any sigmoidal
395 function like $\text{sign}(x)$ and $\tanh(x)$, but in more advanced applications other functions
396 like ReLU [1] are used. The resulting function $f(\vec{x}, \vec{w})$ has then a simple non linear
397 behaviour. It produces a binary output: 1 if the weighted combination is high enough
398 and 0 if it is low enough, with a smooth modulation in-between the two values.

399 The most common architecture for a NN is the so-called *feed-forward* architecture,
400 where many individual perceptrons are arranged in chained layers, which take as input
401 the output of previous layers along with a straight information flux. More complex ar-
402 chitectures could implements also recursive connection, linking a layer to itself, but it
403 should be regarded as sophistication to the standard case. There are endless possibilities

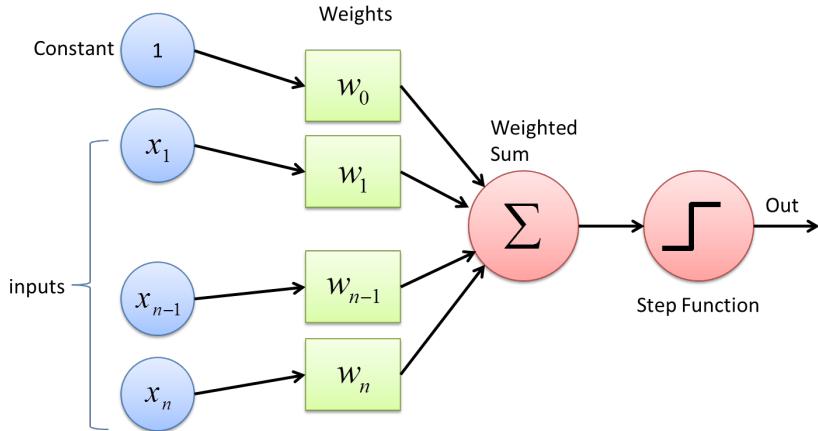


Figure 1.5: Schematic picture of a single layer perceptron. The input vector is linearly combined with the bias factor and sent to an activation function to produce the numerical "binary" output.

of combination and arrangement of neurons inside a NN's layer, but the most simple ones are known as *fully-connected* layers, where every neuron is linked with each other neuron of the following layer, as shown in Figure 1.6. Each connection has its weight, which contributes to modulate the overall combination of signals. The training of a NN consists then in the adjustment and fine-tuning of all the network's weights and parameters through iterative techniques until the desired precision in the output generation is reached.

Although a fully connected network represents the simplest linking choice, the insertion of each weight increases the number of overall parameters, and so the complexity of the model. Thus we want to create links between neurons smartly, rejecting the less useful ones. Depending on the type of data under analysis there are many different established typologies of layers. For example, in the image processing field, the most common choice is the convolutional layer, which implements a sort of discrete convolution on the input data, as shown in Figure 1.7. While processing images, the convolution operation confers to the perception of correlation between adjacent pixels of an image and their color channels, allowing a sort of spatial awareness. Furthermore, the majority of traditional computer vision techniques are based on the discrete convolution of images, and on the features extracted from them.

As a matter of principle a NN with just two successive layers, which is called a *shallow* network, and with an arbitrary number of neurons per layer, can approximate arbitrary well any kind of smooth enough function [32]. However, direct experience suggests that networks with multiple layers, called *deep* networks, can reach equivalent results exploiting a lower number of parameters overall. This is the reason why this discipline goes under the name of *deep* learning: it focuses on deep networks with up

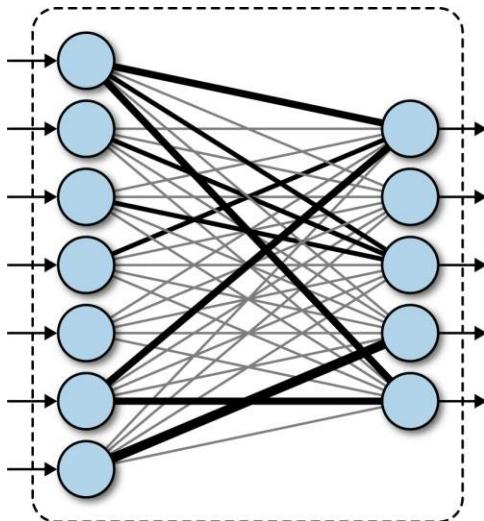


Figure 1.6: Schematic representation of a fully connected (or dense) layer. Every neuron from the first layer is connected with every output neuron. The link thickness represent the absolute value of the combination weight for that particular value.

428 to tens of hidden layers. Such deep structures allow the computation of what is called
 429 deep features, so features of the features of the input data, that allows the network to
 430 easily manage concepts that would be barely understandable for humans.

431 1.3.2 Training Strategies in Deep Learning

432 Depending on the task the NN is designed for, it will have a different architecture and
 433 number of parameters. Those parameters are initialized to completely random values,
 434 tough. The training process is exactly the process of seeking iteratively the right values
 435 to assign to each parameter in the network in order to accomplish the task. The best
 436 start to understanding the training procedure is to look at how a supervised problem is
 437 solved. In supervised problems, we start with a series of examples of true connections
 438 between inputs and correspondent outputs and we try to generalize the rule behind those
 439 examples. After the rule has been picked up the final aim is to exploit it and to apply
 440 it to unknown data, so the new problem could be solved. In opposition to the concept
 441 of supervised problems, there are the *unsupervised* problems, where the algorithm does
 442 not try to learn a rule from a practical example but try to devise it from scratch. A task
 443 typically posed as unsupervised is clustering, when different data are separated in groups
 444 based on the values of their features in the feature space. Usually, only the number of
 445 groups is taken in input from the algorithm, and the subdivision is completely performed
 446 by the machine. In the real world, by the way, there are many different and creative

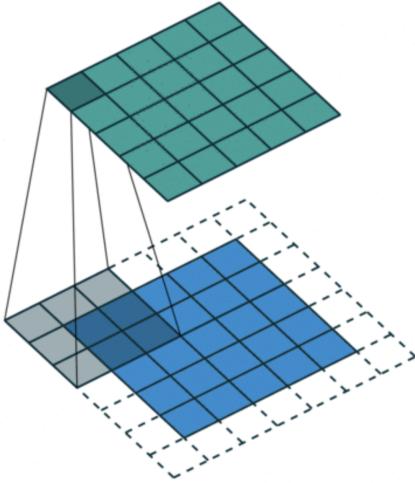


Figure 1.7: Schematic representation of a convolutional layer. The input data are processed by a window kernel that slides all over the image. This operation can recreate almost all the traditional computer vision techniques, and can overcome them, creating new operations, which would be unthinkable to hand-engineered.

447 shapes between pure supervised and pure unsupervised learning, based on the actual
 448 availability of data and specific limitations to the individual task.

449 An interesting mention in this regards should be made about *semi-supervised* learn-
 450 ing, which is typically used in bio-medical applications. This learning technique combines
 451 a large quantity of unlabeled data during training with a limited number of pre-labeled
 452 example. The blend between data can produce considerable improvement in learning
 453 accuracy, and leading to better results respect to pure unsupervised techniques. The
 454 typical situation of usage of this technique is when the acquisition of labeled data re-
 455 quires a highly trained human agent (as an anatomo-pathologist) or a complex physical
 456 experiment. The actual cost of building entire and suitable fully labeled training sets
 457 in these situations would be unbearable, and semi-supervised learnign comes in great
 458 practical help.

459 Another important training technique which worth mentioning is the so-called *cur-
 460 riculum learning* [5]. This is a learning technique inspired by the typical learning curve
 461 human being and animal, that are used to face problems of always increasing difficulty
 462 while learning something new or a new skill. The example presented to the Neural Net-
 463 work are not randomly presented but organized in a meaningful order which illustrates
 464 gradually more concepts, and gradually more complex ones. The experiments show that
 465 through curriculum learning significant improvements in generalization can be achieved.
 466 This approach has both an effect on the speed of convergence of the training process to a

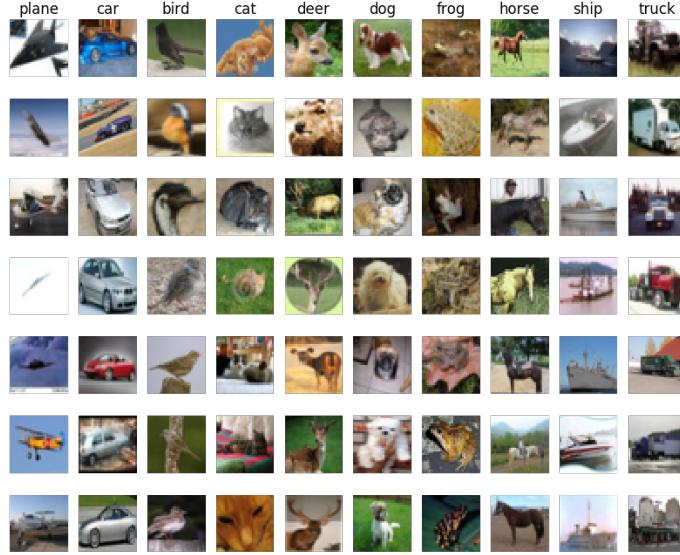


Figure 1.8: Sample grid of images from the CIFAR10 dataset. Each one of the 32×32 image is labeled with one of the ten classes of objects: plane, car, bird, cat, deer, dog, frog, horse, ship, truck.

467 minimum and, in the case of non-convex criteria, on the quality of the local minima ob-
 468 tained. This technique is of particular interest for this work: the generated images, which
 469 will be shown in section 3.3 , will offer a segmentation task much less complex respect to
 470 the analysis of real histological tissue. In the optical of training a DL-based model the
 471 abundantly produced images can be used for the preliminary phase of training, setting
 472 aside the more complex and more valuable hand-labeled images for the finalization of
 473 the training process.

474 1.3.3 Training Algorithms - Error Back-Propagation

475 A good example of supervised problems tough is the classification of images. Let's
 476 assume we have a whole dataset of pictures of different objects (as cats, dogs, cars, etc.)
 477 like the CIFAR10 [22] dataset. This famous dataset is made of over $60K$ labeled colored
 478 images 32×32 divided into 10 categories of objects as shown in Figure 1.8. We could be
 479 interested in the creation of a NN able to assign at every image its belonging class. This
 480 NN could be arbitrarily complex but it certainly will take as input a $32 \times 32 \times 3$ RGB
 481 image and the output will be the predicted class. A typical output for this problem
 482 would be a probability distribution over all the 10 classes like:

$$\vec{p} = (p_1, p_2, \dots, p_{10}), \quad (1.2)$$

$$\sum_{i=1}^{10} p_i = 1, \quad (1.3)$$

and it should be compared with the true label, that is represented just as a binary sequence \vec{t} with the bit correspondent to the belonging class set as 1, and all the others value set to 0:

$$\vec{t} = (0, 0, \dots, 1, \dots, 0, 0). \quad (1.4)$$

Every time an image is given to the model an estimate of the output is produced. Thus, we need to measure the *distance* between that prediction and the true value, to quantify the error made by the algorithm and try to improve the model's predictive power. The functions used for this purpose are called loss functions. The most common choice is the Mean Squared Error (MSE) function that is simply the averaged L^2 norm of the difference vector between \vec{p} and \vec{t} :

$$MSE = \frac{1}{n} \sum_{i=0}^n (t_i - p_i)^2. \quad (1.5)$$

Let's say the NN under training has L consecutive layers, each one with its activation function f^k and its weights vector \vec{w}^k , hence the prediction vector \vec{p} could be seen as the result of the consecutive, nested, application through all the layers:

$$\vec{p} = f^L(\vec{w}^L \cdot (f^{L-1}(\vec{w}^{L-1} \cdot \dots \cdot f^1(\vec{w}^1 \cdot \vec{x}))). \quad (1.6)$$

From both equations 1.5 and 1.6 it is clear that the loss function could be seen as a function of all the weights vectors of every layer of the network. So if we want to reduce the distance between the NN prediction and the true value we need to modify those weights to minimize the loss function. The most established algorithm to do so for a supervised task in a feed-forward network is the so-called *error back-propagation*.

The back-propagation method is an iterative technique that works essentially computing the gradient of the loss function with respect to the weights using the derivative chain rule and updating by a small amount the value of each parameter to lower the overall loss function at each step. Each weight is *moved* counter-gradient, and summing all the contribution to every parameter the loss function approaches its minimum. In equation 1.7 is represented the variation applied to the j^{th} weight in the i^{th} layer in a single step of the method:

$$\Delta w_{ij} = -\eta \frac{\partial E}{\partial w_{ij}}, \quad (1.7)$$

507 where E is the error function, and η is the *learning coefficient*, that modulate the effect of
508 learning through all the training process. This iterative procedure is applied completely
509 to each image in the training set several times, each time the whole dataset is reprocessed
510 is called an *epoch*. The great majority of the dataset is exploited in the training phase
511 to keep running this trial and error process and just a small portion is left out (typically
512 10% of the data) for a final performance test.

513 The loss function shall inevitably be differentiable, and its behavior heavily influences
514 the success of the training. If the loss function presents a gradient landscape rich of
515 local minima the gradient descent process would probably get stuck in one of them.
516 More sophisticated algorithms capable of avoiding this issue have been devised, with
517 the insertion of some degree of randomness in them, as the Stochastic Gradient Descent
518 algorithm, or the wide used *Adam* optimizer [21].

519 While Error-Back Propagation is the most established standard in DL applications, it
520 suffers from some problems. The most common one is the so-called vanishing or exploding
521 gradient issue, which is due to the iterative chain derivation through all the nested
522 level of composition of the function. Without a careful choice of the right activation
523 function and the tuning of the learning hyper-parameters it is very easy to bump into
524 this pitfall. Furthermore the heavy use of derivation rises the inability to handle non-
525 differentiable components and hinders the possibility of parallel computation. However,
526 there are many alternative approaches to network learning beside EBP. The Minimiza-
527 tion with Auxiliary Variables (MAV) method builds upon previously proposed methods
528 that break the nested objective into easier-to-solve local subproblems via inserting aux-
529 iliary variables corresponding to activations in each layer. Such method avoids gradient
530 chain computation and the potential issues associated with it [10]. A further alterna-
531 tive approach to train the network is the Local Error Signals (LES), which is based on
532 layer-wise loss functions. In [29], is shown that layer-wise training can approach the
533 state-of-the-art on a variety of image datasets. It is used a single-layer sub-networks
534 and two different supervised loss functions to generate local error signals for the hidden
535 layers, and it is shown that the combination of these losses helps with optimization in
536 the context of local learning.

537 The training phase is the pulsing heart of a DL model development and it could
538 take even weeks on top-level computers for the most complicated networks. In fact,
539 one of the great limits to the complexity of a network during the designing phase is
540 exactly the available computational power. There are many more further technical details
541 necessary for proper training, the adjustment of which can heavily impact the quality of
542 the algorithm. However, after the training phase, we need to test the performance of the
543 NN. This is usually done running the trained algorithm on never seen before inputs (the
544 test dataset) and comparing the prediction with the ground-truth value. A good way to
545 evaluate the quality of the results is to use the same function used as the loss function
546 during the training, but there is no technical restriction to the choice of this quality
547 metric. The average score on the whole test set is then used as a numerical score for

548 the network, and it allows straightforward comparison with other models' performances,
549 trained for the same task. All this training procedure is coherently customized to every
550 different application, depending on which the problem is posed as supervised or not and
551 depending on the more or less complex network's architecture. The leitmotif is always
552 finding a suitable loss function that quantifies how well the network does what it has
553 been designed to do and trying to minimize it, operating on the parameters that define
554 the network structure.



Figure 1.9: Example of the resulting segmentation mask of an image of an urban landscape. Every interesting object of the image is detected and a solid color region replaces it in the segmentation mask. Every color corresponds to a different class of objects, for example, persons are highlighted in magenta and scooters in blue. The shape and the boundaries of every region should match as precisely as possible the edges of the objects.

555 1.4 Deep Learning-Based Segmentation Algorithms

556 In digital image processing, image segmentation is the process of recognizing and sub-
 557 dividing an image into different regions of pixels that show similar features, like color,
 558 texture, or intensity. Typically, the task of segmentation is to recognize the edges and
 559 boundaries of the different objects in the image and assigning a different label to every
 560 detected region. The result of the segmentation process is an image with the same dimen-
 561 sions of the starting one made of solid color regions, representing the detected objects.
 562 This image is called *segmentation mask*. In Figure 1.9 is shown an example of segmenta-
 563 tion of a picture of an urban landscape: different colors are linked to different classes of
 564 objects like persons in magenta and scooters in blue. This technology has a significant
 565 role in a wide variety of application fields such as scene understanding, medical image
 566 analysis, augmented reality, etc.

567 A relatively easy segmentation problem, and one of the first to be tackled, could
 568 be distinguishing an object from the background in a grey-scale image, like in Figure
 569 1.10. The easiest technique to perform segmentation in this kind of problem is based on
 570 thresholding. Thresholding is a binarization technique based on the image's grey-level
 571 histogram: to every pixel with luminosity above that threshold is assigned the color
 572 *white*, and vice versa the color *black*. However, this is a very primitive and fallacious, yet
 573 very fast method, and it manages poorly complex images or images with un-uniformity
 574 in the background.

575 A lot of other traditional techniques improve this first segmentation method [11].
 576 Some are based on the object's edges recognition, exploiting the sharp change in lumi-
 577 nosity typically in correspondence of the boundary of a shape. Other techniques exploit
 578 instead a region-growing technology, according to which some *seed* region markers are



Figure 1.10: Example of the resulting segmentation mask of an image of a fingerprint obtained through a thresholding algorithm. The result is not extremely good, but this technique is very easy to implement and runs very quickly.

579 scattered on the image, and the regions corresponding to the objects in the image are
 580 grown to incorporate adjacent pixels with similar properties.

581 Every development of traditional computer vision or of Machine Learning-based seg-
 582 mentation algorithm suffers from the same, inevitable limitation. For every one of those
 583 techniques is the designing phase in which the operator should decide precisely which fea-
 584 tures to extract from the image, like different directional derivatives in the image plane
 585 or image entropy, and how to process them for the rest of the analysis. There is thus an
 586 intrinsic limitation in the human comprehension of those quantities and in the possible
 587 way to combine them. The choice is made on the previous experimental results in other
 588 image processing works, and on their theoretical interpretation. Neural Networks instead
 589 are relieved from this limitation, allowing themselves to learn which features are best
 590 suited for the task and how they should be processed during the training phase. The
 591 complexity is then moved on to the design of the DL model and on its learning phase
 592 rather than on the hand-engineering design of the features to extract. In Figure 1.11 an
 593 example high-level feature extracted from a DL model trained for the segmentation of
 594 nuclei in a histological sample. The model learns the typical pattern of arrangement of
 595 nuclei, which would have been impossible to describe equally in advance.

596 1.4.1 State of the Art on Deep Learning Segmentation

597 Similarly to many other traditional tasks, also for segmentation, there has been a thriv-
 598 ing development lead by the diffusion of deep learning, that boosted the performances
 599 resulting in what many regards as a paradigm shift in the field [27].

600 In further detail, image segmentation can be formulated as a classification problem of
 601 pixels with semantic labels (semantic segmentation) or partitioning of individual objects
 602 (instance segmentation). Semantic segmentation performs pixel-level labeling with a set
 603 of object categories (e.g. boat, car, person, tree) for all the pixels in the image, hence it

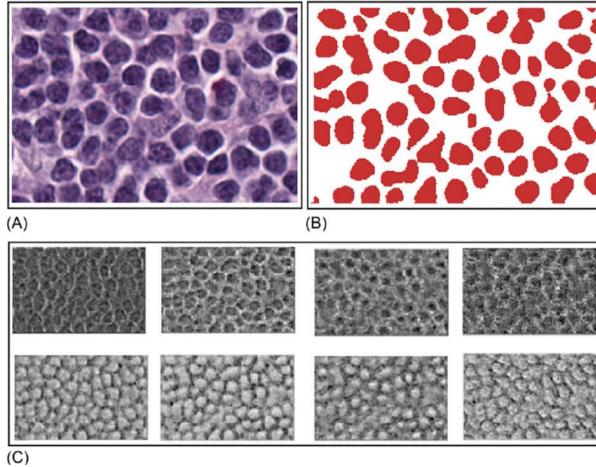


Figure 1.11: (A) An extract from an histological samples, used as input image for the model. (B) The exact segmentation mask. (C) Example of accentuated features during the training: (1-4) for back-ground recognition, (5-8) for nuclei detection. From [2].

is typically a harder task than image classification, which requires just a single label for the whole image. Instance segmentation extends semantic segmentation scope further by detecting and delineating each object of interest in the image (e.g. partitioning of individual nuclei in a histological image).

There are many prominent Neural Network architectures used in the computer vision community nowadays, based on very different ideas such as convolution, recursion, dimensionality reduction, and image generation. This section will provide an overview of the state of the art of this technology and will dwell briefly on the details behind some of those innovative architectures.

Recurrent Neural Networks (RNNs) and the LSTM

The typical application for RNN is processing sequential data, as written text, speech or video clips, or any other kind of time-series signal. In this kind of data, there is a strong dependency between values at a given time/position and values previously processed. Those models try to implement the concept of *memory* weaving connections, outside the main information flow of the network, with the previous NN's input. At each time-stamp, the model collects the input from the current time X_i and the hidden state from the previous step h_{i-1} and outputs a target value and a new hidden state (Figure 1.12). Typically RNN cannot manage easily long-term dependencies in long sequences of signals. There is no theoretical limitation in this direction, but often it arises vanishing (or exploding) gradient problematics during the training phase. A specific type of RNN has been designed to avoid this situation, the so-called Long Short Term Memory (LSTM) [19]. The

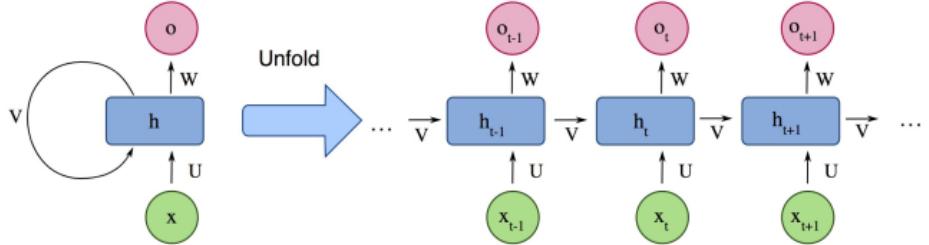


Figure 1.12: Example of the structure of a simple Recurrent Neural Network from [27].

626 LSTM architecture includes three gates (input gate, output gate, forget gate),
 627 which regulate the flow of information into and out from a memory cell, which
 628 stores values over arbitrary time intervals.

629 Encoder-Decoder and Auto-Encoder Models

630 Encoder-Decoder models try to learn the relation between an input and the corre-
 631 sponding output with a two steps process. The first step is the so-called *encoding*
 632 process, in which the input x is compressed in what is called the *latent-space* rep-
 633 resentation $z = f(x)$. The second step is the *decoding* process, where the NN
 634 predicts the output starting from the latent-space representation $y = g(z)$. The
 635 idea underneath this approach is to capture in the latent-space representation the
 636 underlying semantic information of the input that is useful for predicting the out-
 637 put. ED models are widely used in image-to-image problems (where both input
 638 and output are images) and for sequential-data processing (like Natural Language
 639 Processing, NLP). In Figure 1.13 is shown a schematic representation of this ar-
 640 chitecture. Usually, these model follow a supervised training, trying to reduce the
 641 reconstruction loss between the predicted output and the ground-truth output pro-
 642 vided while training. Typical applications for this technology are image-enhancing
 643 techniques like de-noising or super-resolution, where the output image is an im-
 644 proved version of the input image. Or image generation problems (e.g. plausible
 645 new human faces generation) in which all the properties which define the type of
 646 image under analysis should be learned in the representation latent space.

Generative Adversarial Networks (GANs)

The peculiarity of Generative Adversarial Network (GAN) lies in its structure. It is actually made of two distinct and independent modules: a generator and a discriminator, as shown in Figure 1.14. The first module G , responsible for the generation, typically learns to map a prior random distribution of input z to a target distribution y , as similar as possible to the target $G = z \rightarrow y$ (i.e. almost

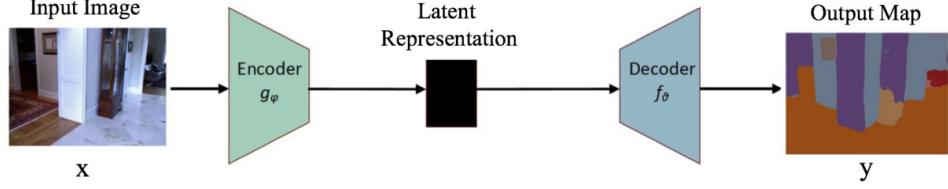


Figure 1.13: Example of the structure of a simple Encoder-Decoder Neural Network from [27].

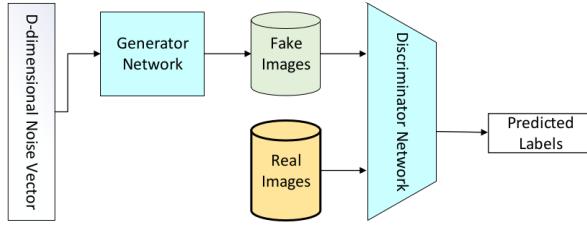


Figure 1.14: Schematical representation of a Generative Adversarial Networks, form [27].

any kind of image-to-image problem could be addressed with GANs, as in [20]). The second module, the discriminator D , instead is trained to distinguish between *real* and *fake* images of the target category. These two networks are trained alternately in the same training process. The generator tries to fool the discriminator and vice versa. The name adversarial is actually due to this *competition* within different parts of the network. The formal manner to set up this adversarial training lies in the accurate choice of a suitable loss function, that will look like:

$$L_{GAN} = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

. The GAN is thus based on a min-max game between G and D . D aims to reduce the classification error in distinguishing fake samples from real ones, and as a consequence maximizing the L_{GAN} . On the other hand, G wants to maximize the D 's error, hence minimizing L_{GAN} . The result of the training process is the trained generator G^* , capable of produce an arbitrary number of new data (images, text, or whatever else):

$$G^* = \arg \min_G \max_D L_{GAN}$$

647 . This peculiar architecture has yielded several interesting results and it has been
 648 developed in many different directions, with influences and contaminations with
 649 other architectures [20].

650 Convolutional Neural Networks (CNNs)

651 As stated before CNNs are a staple choice in image processing DL applications.
652 They mainly consist of three types of layers:

- 653 i convolutional layers, where a kernel window of parameters is convolved with
654 the image pixels and produce numerical features maps.
- 655 ii nonlinear layers, which apply an activation function on feature maps (usually
656 element-wise). This step allows the network to introduce non-linear behavior
657 and then increasing its modeling capabilities.
- 658 iii pooling layers, which replace a small neighborhood of a feature map with some
659 statistical information (mean, max, etc.) about the neighborhood and reduce
660 the spatial resolution.

661 Given the arrangement of successive layers, each unit receives weighted inputs from
662 a small neighborhood, known as the receptive field, of units in the previous layer.
663 The stack of layers allows the NN to perceive different resolutions: the higher-
664 level layers learn features from increasingly wider receptive fields. The leading
665 computational advantage given by CNN architecture lies in the sharing of kernels'
666 weights within a convolutional layer. The result is a significantly smaller number
667 of parameters than fully-connected neural networks. In section 2.7 will be shown a
668 particular application of this architecture, known as *style-transfer* network, which
669 is a particular algorithm capable of implanting the visual texture of a *style* image
670 onto the content of a different image, producing interesting hybrid images. Some
671 of the most notorious CNN architectures include: AlexNet [23], VGGNet [40], and
672 U-Net [34].

673 For this work, U-net architecture is particularly interesting. The U-net model was
674 initially developed for biomedical image segmentation, and in its structure reflects char-
675 acteristics of both CNN and Encoder-Decoder models. Ronneberger et al.[34] proposed
676 this model for segmenting biological microscopy images in 2015. The U-Net architecture
677 is made of two branches, a contracting path to capture context, and a symmetric expand-
678 ing path (see Figure 1.15). The down-sampling flow is made of a Fully Convolutional
679 Network (FCN)-like architecture that computes features with 3×3 kernel convolutions.
680 On the other hand, the up-sampling branch exploits up-convolution operations (or de-
681 convolution), reducing the number of feature maps while increasing their dimensions.
682 Another characteristic of this architecture is the presence of direct connections between
683 layers of a similar level of compression in compressing and decompressing branches.
684 Those links allow the NN to preserve spatial and pattern information. The Network
685 flow eventually ends with a 1×1 convolution layer responsible for the generation of the
686 segmentation mask of the input image.

687 A recent example of a practical application of a CNN to histological images could
688 be found in [28]. In this work the Inception v3 is trained on hand-labeled samples of

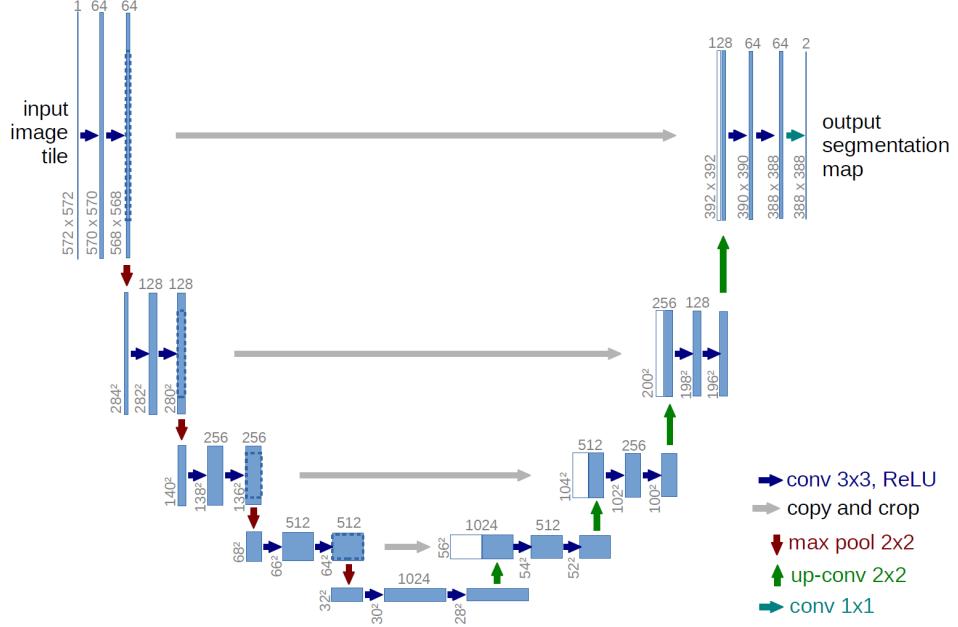


Figure 1.15: Scheme of the typical architecture of a U-net NN. This particular model was firstly proposed by Ronneberger et al. in [34].

pancreatic tissue (like in Figure 1) to recognise tumoral regions from healthy ones in a pancreatic tissue specimen treated with Ki67 staining. The Inception v3 [33] network is a deep convolutional network developed by Google, trained for object detection and image classification on the ImageNet dataset [15]. Recognition of tumoral region of Ki67 stained pancreatic tissue samples is based on the detection and counting of some specific marker cells. In Figure 1.16 is shown a pair of original image and the computed segmentation mask, which label in red tumoral regions and in green the healthy ones. This work is based on a technique called **transfer learning**, which consists in the customization and specialization of a pre-trained NN previously trained for similar, but essentially different, tasks. The final part of the training of this version of Inception v3 has been performed of a dataset of 33 whole slide images of Ki67 stained neuroendocrine tumor biopsies acquired from 33 different patients, digitized with a $20\times$ magnification factor and successively divided in 64×64 patches.

1.4.2 Image Segmentation Datasets

Besides the choice of suitable architecture the most important aspect while developing a NN is the dataset on which perform the training process. Let's confine the discussion only to image-to-image problems, like segmentation problems. There are a lot of widely used datasets, but I want to mention just a few of them to give the idea of their typical

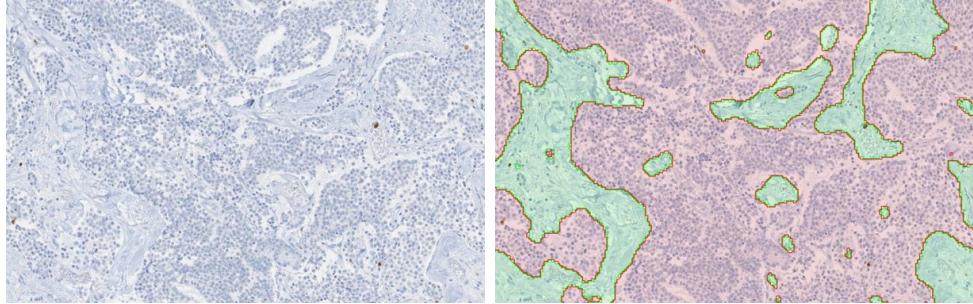


Figure 1.16: (left) Original image of a Ki67 stained pancreatic tissue sample, (right) the corresponding segmentation mask, which label in red tumoral regions and in green the healthy ones. From [28].

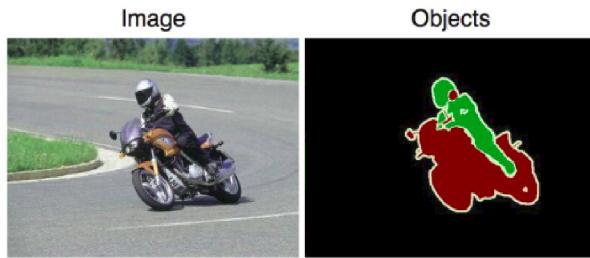


Figure 1.17: An example image from the PASCAL dataset and its corresponding segmentation mask [17].

707 characteristics.

708 A good example of segmentation is the Cityscapes dataset [12], which is a large-scale
 709 database with a focus on semantic understanding of urban street scenes. The dataset
 710 is made of video sequences from the point of view of a car in the road traffic, from 50
 711 different cities in the world. The clips are made of 5K frames, labeled with extremely high
 712 quality at pixel-level and an additional set of 20K weakly-annotated frames. Each pixel
 713 in the segmentation mask contains the semantic classification, among over 30 classes of
 714 objects. An example of an image from this dataset is shown in Figure 1.9.

715 The PASCAL Visual Object Classes (VOC) [17] is another of the most popular
 716 datasets in computer vision. This dataset is designed to support the training of algo-
 717 rithms for 5 different tasks: segmentation, classification, detection, person layout, and
 718 action recognition. In particular, for segmentation, there are over 20 classes of labeled
 719 objects (e.g. planes, bus, car, sofa, TV, dogs, person, etc.). The dataset comes divided
 720 into two portions: training and validation, with 1,464 and 1,449 images, respectively. In
 721 Figure 1.17 is shown an example of an image and its corresponding segmentation mask.

722 As last mention I would report t he ImageNet project [15], which is a large visual
 723 database designed for use in visual object recognition software research. It consists

724 of more than 14 million images have been hand-annotated by the project to indicate
725 what objects are pictured and in at least one million of the images, bounding boxes
726 are also provided. ImageNet contains more than 20,000 categories of objects. Since
727 2010, the ImageNet project runs an annual software contest, the ImageNet Large Scale
728 Visual Recognition Challenge (ILSVRC), where software programs compete to correctly
729 classify and detect objects and scenes. This kind of competitions is very important
730 for the research field, as it inspire and encourage the development of new models and
731 architectures.

732 It is worth mentioning that in the medical image processing domain typically the
733 available dataset is definitely not that rich and vast (that is actually the seed of this
734 work) and thus many techniques of data augmentation have been devised, to get the
735 best out of the restricted amount of material. Generally, data augmentation manipu-
736 lates the starting material applying a set of transformation to create new material, like
737 rotation, reflection, scaling, cropping and shifting, etc. Data augmentation has been
738 proven to improve the efficacy of the training, making the model less prone to over-
739 fitting, increasing the generalization power of the model, and helping the convergence to
740 a stable solution during the training process.

⁷⁴¹ Chapter 2

⁷⁴² Technical Tools for Model ⁷⁴³ Development

⁷⁴⁴ As mentioned in the introduction, this project wants to produce synthetic histological
⁷⁴⁵ images paired with their corresponding segmentation mask, to train Neural Networks
⁷⁴⁶ for the automatization of real histological images analysis. The production of artificial
⁷⁴⁷ images passes through the processing of a three dimensional, virtual model of a his-
⁷⁴⁸ tological structure, which is the heart of this thesis work. The detailed description of
⁷⁴⁹ the development of the two proposed histological models will follow the present chapter
⁷⁵⁰ and will occupy all the chapter 3. Here I will dwell, instead, on every less common
⁷⁵¹ tool employed during the models' designing phase. From the practical point of view, this
⁷⁵² project is quite articulated and the development has required the harmonization of many
⁷⁵³ different technologies, tools, and code libraries. The current chapter should be seen as a
⁷⁵⁴ theoretical complement for chapter 3, and its reading is suggested to the reader for any
⁷⁵⁵ theoretical gap or for any further technical deepening. The reader already familiar with
⁷⁵⁶ those technical tools should freely jump to the models' description.

⁷⁵⁷ All the code necessary for the work has been written in a pure Python environment,
⁷⁵⁸ using several already established libraries and writing by my self the missing code for
⁷⁵⁹ some specific applications. I decided to code in Python given the thriving variety of
⁷⁶⁰ available libraries geared toward scientific computation, image processing, data analysis,
⁷⁶¹ and last but not least for its ease of use (compared to other programming languages).
⁷⁶² In each one of the following subsections I will mention the specific code libraries which
⁷⁶³ have been employed in this project for every technical necessity.

⁷⁶⁴ 2.1 Quaternions

⁷⁶⁵ Quaternions are, in mathematics, a number system that expands to four dimensions the
⁷⁶⁶ complex numbers. They have been described for the first time by the famous mathemati-

767 cian William Rowan Hamilton in 1843. This number system define three independent
768 *imaginary* units $\mathbf{i}, \mathbf{j}, \mathbf{k}$ as in (2.1), which allows the general representation of a quaternion
769 \mathbf{q} is (2.2) and its inverse \mathbf{q}^{-1} (2.3) where a, b, c, d are real numbers:

$$\mathbf{i}^2 = \mathbf{j}^2 = \mathbf{k}^2 = \mathbf{ijk} = -1, \quad (2.1)$$

$$\mathbf{q} = a + bi + cj + dk, \quad (2.2)$$

$$\mathbf{q}^{-1} = (a + bi + cj + dk)^{-1} = \frac{1}{a^2 + b^2 + c^2 + d^2} (a - bi - cj - dk). \quad (2.3)$$

770 Furthermore, the multiplication operation between quaternionn does not benefit from
771 commutativity, hence the product between basis elements will behave as follows:

$$\begin{aligned} \mathbf{i} \cdot 1 &= 1 \cdot \mathbf{i} = \mathbf{i}, & \mathbf{j} \cdot 1 &= 1 \cdot \mathbf{j} = \mathbf{j}, & \mathbf{k} \cdot 1 &= 1 \cdot \mathbf{k} = \mathbf{k} \\ \mathbf{i} \cdot \mathbf{j} &= \mathbf{k}, & \mathbf{j} \cdot \mathbf{i} &= -\mathbf{k} \\ \mathbf{k} \cdot \mathbf{i} &= \mathbf{j}, & \mathbf{i} \cdot \mathbf{k} &= -\mathbf{j} \\ \mathbf{j} \cdot \mathbf{k} &= \mathbf{i}, & \mathbf{k} \cdot \mathbf{j} &= -\mathbf{i}. \end{aligned} \quad (2.4)$$

772 This number system has plenty of peculiar properties and applications, but for this
773 project, quaternions are important for their ability to represent, in a very convenient way,
774 rotations in three dimensions. The particular subset of quaternions with vanishing real
775 part ($a = 0$) has a useful, yet redundant, correspondence with the group of rotations in
776 tridimensional space $\text{SO}(3)$. Every 3D rotation of an object can be represented by a 3D
777 vector \vec{u} : the vector's direction indicates the axis of rotation and the vector magnitude $|\vec{u}|$
778 express the angular extent of rotation. However, the matrix operation which expresses
779 the rotation around an arbitrary vector \vec{u} it is quite complex and does not scale easily
780 for multiple rotations [7], which brings to very heavy and entangled computations.

781 Using quaternions for expressing rotations in space, instead, it is very convinient.
782 Given the unit rotation vector \vec{u} and the rotation angle θ , the corresponding rotation
783 quaternion \mathbf{q} becomes (2.6):

$$\vec{u} = (u_x, u_y, u_z) = u_x \mathbf{i} + u_y \mathbf{j} + u_z \mathbf{k}, \quad (2.5)$$

$$\mathbf{q} = e^{\frac{\theta}{2}(u_x \mathbf{i} + u_y \mathbf{j} + u_z \mathbf{k})} = \cos \frac{\theta}{2} + (u_x \mathbf{i} + u_y \mathbf{j} + u_z \mathbf{k}) \sin \frac{\theta}{2}, \quad (2.6)$$

$$\mathbf{q}^{-1} = \cos \frac{\theta}{2} - (u_x \mathbf{i} + u_y \mathbf{j} + u_z \mathbf{k}) \sin \frac{\theta}{2}, \quad (2.7)$$

784 where in (2.6) we can clearly see a generalization of the Euler's formula for the
785 exponential notation of complex numbers, which hold for quaternions. It can be shown

786 that the application of the rotation represented by \mathbf{q} on an arbitrary 3D vector \vec{v} should
787 be easily expressed as:

$$\vec{v}' = \mathbf{q}\vec{v}\mathbf{q}^{-1}, \quad (2.8)$$

788 using the Hamilton product defined on quaternions (2.4). This rule raises a very con-
789 vinient and an extremely scalable way to compute consecutive rotations in space. Given
790 two independent and consecutive rotations represented by the two quaternions \mathbf{q} and \mathbf{p}
791 applied on the vector \vec{v} the resulting rotated vector \vec{v}' is simply yielded as:

$$\vec{v}' = \mathbf{p}(\mathbf{q}\vec{v}\mathbf{q}^{-1})\mathbf{p}^{-1} = (\mathbf{pq})\vec{v}(\mathbf{qp})^{-1}, \quad (2.9)$$

792 which essentially is the application of the rotation $\mathbf{r} = \mathbf{qp}$ on the vector \vec{v} . This repre-
793 sentation is completely coherent with the algebra of 3D rotations, which does not benefit
794 from commutativity in turn.

795 Given this property, quaternions are indeed widely used in all sorts of applications
796 of digital 3D space design, as for simulations and videogame. The position of an object
797 in the space in simulations is generally given by the application of several independent
798 rotations, typically in the order of a tenth of rotations, which with quaternions is given
799 easily by the product of simple objects. Every other alternative method would imply the
800 use of matrix representation of rotations or other rotation systems as Euler's angles and
801 would eventually make the computation prohibitive.

802 The use of quaternions in this work will be justified in section 3.1, while speaking of
803 parametric L-systems in 3D space, used to build the backbone of the ramified structure
804 of blood vessels in the reconstruction of a sample of pancreatic tissue.

805 I was able to find many Python libraries for computation with quaternions, but the
806 one I appreciated the most for its interface and ease of use was the `pyquaternion`. With
807 this library, it's immediate the definition of a quaternion by its correspondent rotation
808 vector, and the multiplication between quaternions is straightforward.

809 2.2 Parametric L-Systems

810 Lindenmayer systems, or simply L-systems, were conceived as a mathematical theory of
811 plant development [25] in 1968 by Aristid Lindenmayer. Successively, a lot of geometri-
812 cal interpretations of L-systems were proposed to make them a versatile instrument for
813 modeling the morphology typical of plants and other organic structures. As a biologist,
814 Lindenmayer studied different species of yeast and fungi and worked the growth patterns
815 of various types of bacteria (e.g. as the *cyanobacteria Anabaena catenula*). The main
816 purpose for which L-systems were devised was to allow a formal description of the devel-
817 opment of simple multicellular living organisms. Subsequently, the potentiality of these
818 systems was expanded to describe higher-order plants and complex branching structures.

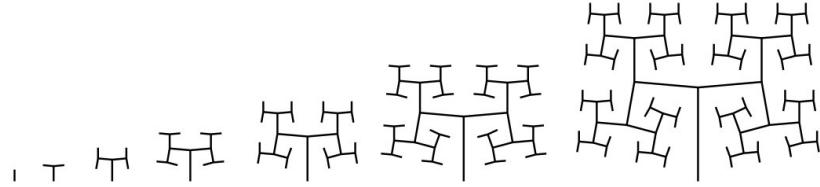


Figure 2.1: Growth pattern for the space-filling fractal-like system, used to mimic the blood vessel bifurcations in sec 3.1.

819 An L-system is in general defined by an *axiom* sequence and some development *rules*,
 820 which are recursively applied to the sequence and lead its development. The original
 821 proposed L-system was fairly simple and shows really well the idea underneath:

$$\begin{aligned} \textit{axiom} &: A \\ \textit{rules} &: (A \rightarrow AB), \quad (B \rightarrow A) \end{aligned}$$

822 where A and B could be any two different patterns in the morphology of an algae, or
 823 could be different bifurcations in a ramified structure. The iterative application of the
 824 rules to the axiom sequence, let's say for 7 times, will produce the following sequence:

$$\begin{aligned} n = 0 & : A \\ n = 1 & : AB \\ n = 2 & : ABA \\ n = 3 & : ABAAB \\ n = 4 & : ABAABABA \\ n = 5 & : ABAABABAABAAB \\ n = 6 & : ABAABABAABAABABAABABA \\ n = 7 & : ABAABABAABAABABAABAABABAABAAB . \end{aligned}$$

825 This kind of tool, as will be shown also in 3.1, is particularly suited for the creation of
 826 structures with fractal behavior, and it has been used in this work to create the backbone
 827 of the entangled bifurcation in blood vessels in the modelization of pancreatic tissue. In
 828 particular, there was the need for a fractal-like space-filling ramification as the one shown
 829 in Figures 2.1.

830 The system in Figure 2.1 represent the successive ramification of a structure which
 831 grows adding segments gradually shorter, by a lenght ratio parameter R and inclined
 832 of $\delta = \pm 85^\circ$ respect the previous branch. The axiom and the rules that produce this
 833 structure are the following:

$$\begin{aligned}
& axiom : A & (2.10) \\
& rule_1 : A \rightarrow F(1)[+A][-A] \\
& rule_2 : F(s) \rightarrow F(s \cdot R)
\end{aligned}$$

where A represent the start of a new branch and $F(s)$ represent a branch of lenght s .
 The presence of a rule which acts differently depending on the target object, is an further
 sophistication respect to the standard L-system. For this reason these systems are called
parametric L-systems.

The use of standard L-systems turned out to be widespread, and there were a lot of
 different Python libraries at my disposal for coding. By the way, parametric L-systems
 were not just as popular, and I was not able to find a reliable library on which to build
 my work. I decided then to code a parametric branching system able to recreate the
 structure with rules (2.10) at any desired level of iteration. Having created the tool I
 needed on my own I was able to add all the optional features I would have needed during
 the development, like an adjustable degree of angular noise in the branch generation.

2.3 Voronoi Tassellation

Voronoi diagrams, or Voronoi decompositions, are space-partitioning systems, which
 divides an n -dimensional Euclidian space into sub-regions depending on the proximity
 to a given set of objects. More precisely, given an n -dimensional space and m starting
 point p_1, \dots, p_m inside it, the whole space will be subdivided in m adjacent regions. Every
 point of the space is assigned to the region correspondent to the nearest starting point.
 In Figure 2.2 is shown a practical example of a Voronoi decomposition of a plane into
 20 regions corresponding to the 20 starting points. Informal use of Voronoi diagrams
 can be traced back to Descartes in 1644, and many other mathematicians after him.
 But, Voronoi diagrams are named after Georgy Feodosievych Voronoy who defined and
 studied the general n-dimensional case in 1908 [46].

More precisely, let X be a metric space and d the distance defined on it. Let K be
 the set of indices and let $(P_k)_{k \in K}$ be the tuple of sites in the space X . The k^{th} Voronoi
 cell R_k , associated with the site P_k is the set of all the points in X whose distance to P_k
 is smaller than the distance to any other site P_j , with $j \neq k$, or in other words:

$$R_k = \{x \in X \mid d(x, P_k) \leq d(x, P_j) \forall j \in K, j \neq k\}, \quad (2.11)$$

depending on the notion of distance defined on the space X the final redistribution in
 subregions will look very differently.

In addition to the choice of the distance function, another fundamental factor is the
 distribution of sites in the space to be divided. If the points are chosen equally and ho-
 mogenously distributed the final distribution will appear as a simple regular lattice, while

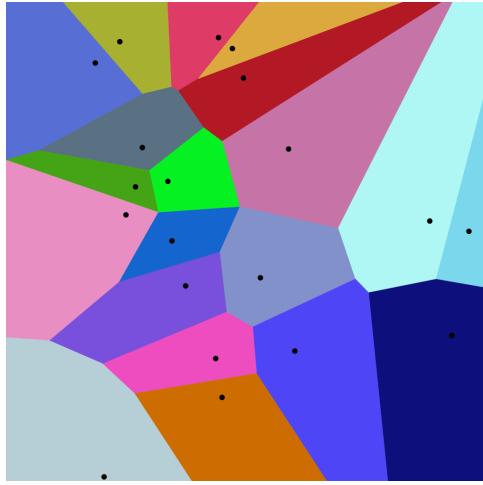


Figure 2.2: Example of a Voronoi decomposition of a plane into 20 regions corresponding to 20 starting points.

a completely random distribution of points in the space will provide a decomposition in cells with very different shapes and volumes, as shown in Figure 2.3. Interesting results concerning points from a semi-random distribution will be shown in section 3.1, which leads to decomposition with a good richness in shapes but with the desired homogeneity in volumes.

The Voronoi decomposition has been of great interest in this project for the division of a 3D space in subregions, to recreate the spatial distribution of cells in a sample of human tissue, as will be shown in section 3.1. The formal definition of Voronoi regions (2.11) ensures the convexity of each decomposition's tassel, which in three-dimensional space would be adjacent convex polyhedrons. Every tassel of the decomposition will be represented by a bounded 3-dimensional convex hull¹, with except for those most external cells which are unbounded and requires special attention while using.

The most widespread tool for the computation of Voronoi decompositions in Python is contained in the `spatial` submodule of the famous library `SciPy` [45], which is a staple tool for an incredible variety of scientific algorithms. The `Voronoi` object from `Scipy` library offers a very efficient algorithm for space-partitioning, and it has been one of the pillars for the modelization of tissues. Unluckily this module does not easily allow to perform Voronoi decomposition with different definitions of distance functions d other than the Euclidian distance, which would have allowed interesting studies.

¹See section 2.5 for further details.

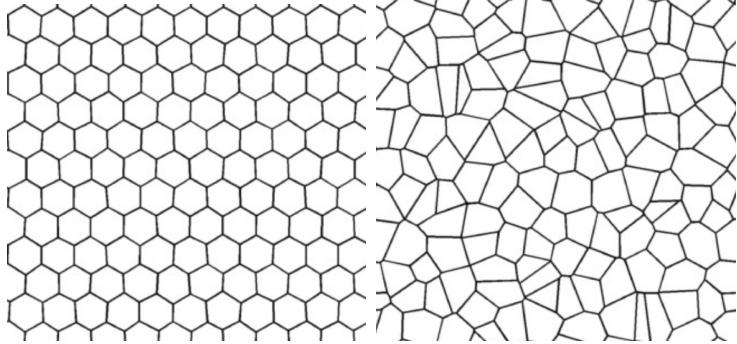


Figure 2.3: On the right an example of 2D Voronoi decomposition resulting from homogeneously distributed points in the plane. On the left the resulting decomposition obtained from randomly distributed points in the plane, from [3].

884 2.4 Saltelli Algorithm - Random Number Generation

885 As mentioned in section 2.3, in this project there was the need for quasi-random number
 886 generation for the production of Voronoi tessellations. Quasi-random sequences (or low-
 887 discrepancy sequences) are patterns of numbers that emulate the behavior of uniform
 888 random distributions but have a more homogeneous and quick coverage of the sampling
 889 domain, which provides an important advantage in applications as in quasi-Monte Carlo
 890 integration techniques, as shown in Figure 2.4. In computer science there is not any
 891 possibility of recreating *true* random sequences, hence any stochasticity is completely
 892 deterministic in its essence even if produced by very chaotic processes². Indeed, every
 893 algorithm for random number generation is completely repeatable given its starting sta-
 894 tus. Quasi-random sequences are completely deterministic too, but implements more
 895 *regular*, well-behaved algorithms.

896 A first good example to understand the concept of quasi-random generation could be
 897 an additive recurrence, as the following:

$$s_{n+1} = (s_n + \alpha) \bmod 1, \quad (2.12)$$

898 which for every seed element s_0 and real parameter α produced completely different
 899 sequences.

900 In the bottom line of Figure 2.4 is clearly visible the good and homogeneous coverage
 901 of the sampling domain, although it is strongly visible a regular pattern between points,
 902 which does not convey an *organic* sensation at all. However, increasing the complexity of

²A chaotic process is a deterministic process which has an extremely sensible dependence on its starting conditions. This property mimics very effectively the behavior of true random processes, which are intrinsically forbidden in computer science.

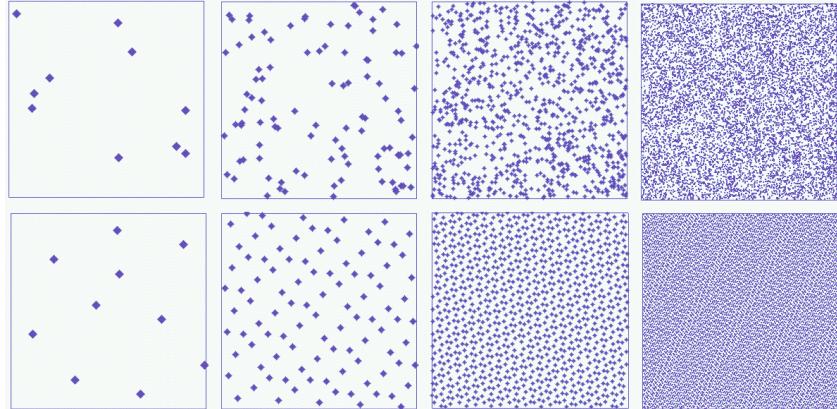


Figure 2.4: Coverage of the unit square with an additive quasirandom numbers sequence as in 2.12 (up) and for uniformly sampled random numbers (bottom). From left to right: 10, 100, 1000, 10000 points.

903 our very simple starting model 2.12 it is possible to overcome this *artificial* appearance
 904 of sampled points and to produce very good samples.

905 A notorious algorithm for quasi-random number generation is the Sobol sequence,
 906 introduced by the russian mathematician Ilya M. Sobol in 1967 [43]. In its work, Sobol
 907 wanted to construct a sequence x_n of points in the s -dimensional unitary hypercube
 908 $I^s = [0, 1]^s$ such as for any integrable function f :

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n f(x_i) = \int_{I_s} f. \quad (2.13)$$

909 Sobol wanted to minimize the *holes* in the sampled domain (which it could be shown
 910 to be a property that helps the convergence of the sequence) and minimize as well
 911 the *holes* in every lower-dimension projection of the sampled points. The particularly
 912 good distributions that fulfill those requirements are known as (t, m, s) -nets and (t, s) -
 913 sequences in base b .

914 To better understand them we need first to define the concept of s -interval in base b ,
 915 which is a subset of I_s such as:

$$E_s^b = \prod_{j=1}^s \left[\frac{a_j}{b^{d_j}}, \frac{a_j + 1}{b^{d_j}} \right), \quad (2.14)$$

916 where a_j and d_j are non-negative integers, and $a_j < b^{d_j}$ for all j in $\{1, \dots, s\}$.

917 Let be t and m two integers such as $0 \leq t \leq m$. A (t, m, s) -net in base b is defined
 918 as a sequence x_n of b^m points of I_s such that:

$$\text{Card } \mathbf{P} \cap \{x_1, \dots, x_n\} = b^t \quad (2.15)$$

919 for all the elementary interval \mathbf{P} in base b of hypervolume $\lambda(\mathbf{P}) = b^{t-m}$.

920 Given a non-negative integer t , a (t, s) -sequence in base b is an infinite sequence of
921 points x_n such that for all integers $k \geq 0$, $m \geq t$ the sequence $\{x_{kb^m}, \dots, x_{(k+1)b^m-1}\}$ is
922 a (t, m, s) -net in base b .

923 Sobol in his article described in particular (t, m, s) -net and (t, s) -sequence in base 2.
924 A more thorough description of all the formal properties of those particular sequences
925 could be found in [42].

926 In order to perform the actual sampling during the modelization, it has been used the
927 `saltelli` module from the `SALib` library, which performs sampling in an s -dimensional
928 space following the Saltelli algorithm, which is a specific improved version of the Sobol
929 algorithms oriented toward the parameter sensitivity analysis [36], [37].

930 2.5 Planar Section of a Polyhedron

931 As will be shown in section 3.1 a fundamental step for the functioning of the modelization
932 is the planar section of a three-dimensional polyhedron. It turned out that there is
933 no general rule to perform a planar section of a convex polyhedron with an arbitrary
934 number of faces, respect to an arbitrary sectioning plane. Hence, I devised an algorithm
935 to handle this task. In the case of a full intersection, the result of the sectioning process
936 of a polyhedron is a polygonal surface, otherwise, it could be an empty set of points or
937 a segment in case of particular tangency, but those two cases are not of interest to the
938 model. This tool has been created from scratch by me in the preliminar design phase and
939 also the implemented algorithm has been devised by me and chosen as the best option
940 among other possibilities. It is not fruit of an extended and thorough formal analysis,
941 hence it has not the pretece to be an extremely optimized algortihm. Nevertheless, this
942 tool passed all the test I required and yielded the expected results.

943 Given a convex polyhedron with n vertices and a sectioning plane p , let V be the set
944 of all the vertices and $f_p(\vec{x})$ the equation defining the plane. The algorithm is defined
945 by the following steps:

- 946 1. Divide V in two subsets: A made of those vertices which lie above and B , made of
those which lie below the sectioning plane. Like in 2.16:

$$A = \{\vec{v} \in V \mid f_p(\vec{v}) \geq 0\} \quad (2.16)$$
$$B = \{\vec{v} \in V \mid f_p(\vec{v}) \leq 0\}$$

946 If any of the two subsets tourns out to be empty the plane p does not intersect the
947 polyhedron, and the section is empty. A and B are represented in different colors
948 in Figure 2.5.

- 949 2. Detect, and *draw*, any possible line that crosses two points respectively from A and
 B . If n_A and n_B are the numbers of points above and below the plane then there

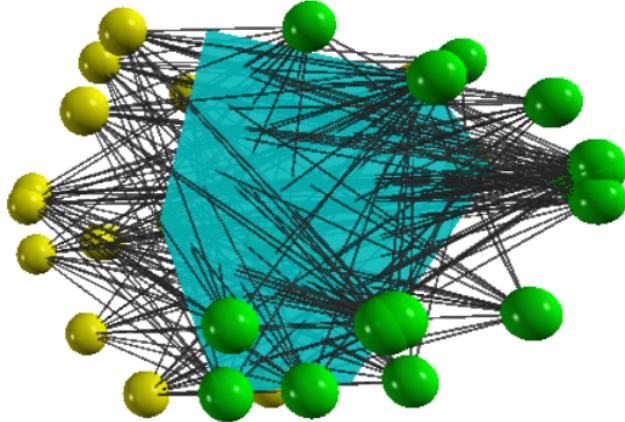


Figure 2.5: In this picture is shown an example of the application of the algorithm for the planar section on a polyhedron. The vertices are divided into two groups, with different colors yellow and green. All the possible lines between any couple of vertices picked from the two classes are drawn in dark gray. In Turquoise the resulting planar section, obtained as the convex hull containing all the intersections between the lines and the plane.

951 will be $n_A \times n_B$ possible lines. In Figure 2.5 all the lines between the two classes
 952 of points are drawn in dark gray.

- 953 3. Detect P , the set of all the points from the intersection between the $n_A \times n_B$ lines
 954 from the previous step and the sectioning plane p . All these points will lie on the
 955 same plane, within the boundaries of the polygonal section.
- 956 4. The final polygon is then yielded by computing the convex hull of the points in
 957 P . The convexity of the starting polyhedron in fact ensures the convexity of any
 958 section of the solid.

959 The result of the algorithm is, as just stated, a convex hull, which in geometry
 960 is defined as the smallest convex envelope or convex closure of a set of points. In 2
 961 dimensions is the smallest convex polygon containing a certain set of points in a plane
 962 (In Figure 2.6 is shown the so-called “rubber band effect”), and in 3 dimensions it is the
 963 smallest convex polyhedron containing a set of points in the space.

964 In Python, the most convenient way to work with convex hulls was to use the sub-
 965 module `spatial.ConvexHull` from the `SciPy` library [45]. This module allows also a
 966 convenient way for plotting images with `Matplotlib`, which is the point of reference for
 967 plotting and image formation in Python.

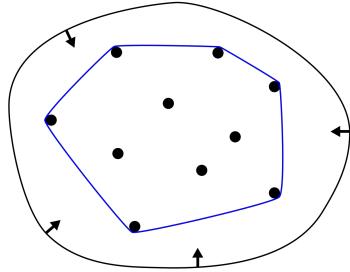


Figure 2.6: Representation of the convex hull of a bounded planar set of points. This particular enclosure goes under the name of "rubber band effect".

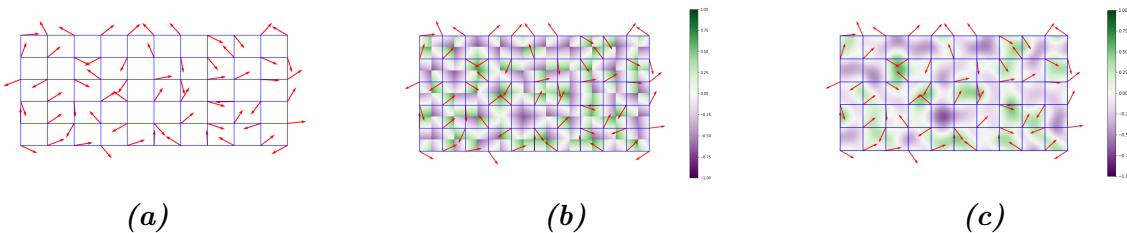


Figure 2.7: The three main steps of the algorithm to produce Perlin noise. In 2.7a the plane discretization and the assignment of a gradient vector to every node of the grid. In 2.7b the computation of the dot product with all the points inside the discretization and in 2.7c the interpolation of the values to create the final function.

968 2.6 Perlin Noise

969 Perlin noise is a widely used form of noise in computer graphics, which mimics very
 970 well natural and smooth fluctuations around a constant value. It has been developed by
 971 Ken Perlin in 1983, and it is now the staple tool for giving texture to object in virtual
 972 modelization, often considered the *salt* of computer graphics texturization. The Perlin
 973 noise is a gradient-based algorithm defined on grid discretization of a n -dimensional
 974 space. The algorithm involves three subsequent steps:

- 975 1. The first step is to discretize the n -dimensional space in a regular lattice: the
 976 dimension of the grid will impact heavily on the scale of the noise. As in Figure
 977 2.7a at every node of the grid is assigned a randomly oriented n -dimensional unitary
 978 gradient vector. This is the preliminary setup which will allow the computation of
 979 the actual noise function in every point of the space.
- 980 2. Given the candidate point \vec{x} between the grid nodes onto which evaluate the noise
 981 there are 2^n nearest grid nodes. For each one of these 2^n nodes, it is evaluated the

982 distance vector from \vec{x} as the offset between the two points. Then it is computed
983 the dot product between every pair made of nearby gradient vectors and the offset
984 vector. This operation should be thought of as made on every point in the lattice,
985 as in Figure 2.7b, where at every point of the grid is represented just one of the
986 $2^2 = 4$ series of dot products.

- 987 3. The final step is the interpolation between the 2^n series of dot products. To per-
988 form the interpolation usually is used a function with vanishing first degree (and
989 preferably also second degree) derivative in correspondence of the 2^n grid nodes ³.
990 This means that the noise function will pass through zero at every node and have
991 a gradient equal to the pre-computed grid node gradient. These properties give
992 Perlin noise its characteristic spatial scale and smoothness.

993 In general, the final result of the algorithm is a smooth function with a random-
994 like behavior that mimics really well an organic appearance, like in Figure 2.8, with
995 fluctuation around the value 0, with amplitude $\in [0; 1]$. The surface in Figure 2.8 has
996 been produced plotting in 3D the results of the function `pnoise2` from the library `noise`,
997 which offers a tool for the production of different type of noise. In order to put in practical
998 use this module some adjustment were required. The particular function `pnoise2` simply
999 yields the value of the Perlin noise surface in correspondence to a single (x, y) point in the
1000 plane in a deterministic way. There was not the possibility to generate different whole
1001 Perlin surfaces every run. To overcome this limitation I made a vectorized⁴ version of
1002 `pnoise2` able to evaluate the function over an arbitrary wide grid of points expressed as
1003 set of all the pairs of coordinates (x, y) of the grid's nodes. In this way a single call to
1004 the function is able to produce the entire surface covering the grid, in the form of a single
1005 NumPy 3D-array. Furthermore, to recover the possibility of generating always different
1006 surfaces as in a random generation, I inserted an offset coordinate (x_O, y_O) , which moves
1007 in the plane the origin of the surface generation. This pair of offset coordinates then acts
1008 also as a *seed* in the generation, allowing to completely recreate previously generated
1009 material in a controlled way.

1010 2.7 Style-Transfer Neural Network

1011 Style-Transfer Neural Networks are common models, able of creating new hybrid images
1012 implanting the visual style from an image preserving the visual content of another image.

³Usually are used functions with a sigmoidal behavior, like any smoothstep function, which is a family of very common items in computer graphics.

⁴In Python the staple tool for scientific, number-cruncher computation is the NumPy library, which allows a fast, complete and efficient way to perform computation between number structure. The operation of transforming a function which acts on a single value (or pair of values) to a function able to perform on a suitable datastructure is called *vectorization*, and its the recomended way to proceed when handling numerical functions in Python.

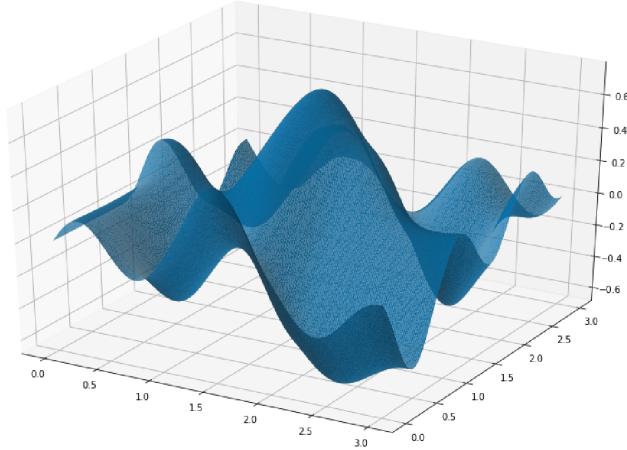


Figure 2.8: Example of Perlin noise 2D function produced while working on this project. This surface offers a smooth variation around the value 0 with amplitude $\in [0; 1]$.

1013 The two images necessary for the algorithm are called *style* image S and *content* image
 1014 C , and the resulting *styled* picture X , as in Figure 2.9.

1015 There are many different tested and comparable architectures to compute this kind
 1016 of algorithm. In my work I decided to use in particular the procedure described in [18],
 1017 using the PyTorch ecosystem to implement the necessary code.

1018 The backbone of the architecture is the VGG-19 network, which is a convolutional
 1019 neural network 19 layers deep, as in Figure 2.10. This huge model has been pre-trained
 1020 on over a million images from the ImageNet database [16], for the classification into
 1021 over than 1000 classes of objects. As a result, the network has learned rich feature
 1022 representations for a wide range of images. The best (and conceptually the only) way to
 1023 load a pre-trained model is to load the ordered set of weights that define the network and
 1024 to initialize an empty module with those values. This is the perfect start for creating a
 1025 style transfer network, which requires a further and briefer training phase, to completely
 1026 customize the network.

1027 The key ingredient for finalizing the model is to insert some little but fundamental
 1028 modifications and to extend the training on the pair of input images. This final training
 1029 should be aware of the *concepts* of the visual style and visual content of the image, and
 1030 the operation should try to preserve them both. This is usually done minimizing two
 1031 new loss function, computed between the staring image and the produced image:

1032 Content Loss

1033 The content loss is a function that represents a weighted version of the content
 1034 distance for an individual layer. The most commonly used function to evaluate the
 1035 preservation of content between two images is the simple Mean Squared Error as



Figure 2.9: Different examples of application of a style-transfer NN on the same content image, with different style images, from [18]. The original picture depicts the Neckarfront in Tbingen, Germany (TOP-LEFT). The painting used as style image shown in the bottom left corner of each panel are in clockwise order: • The Shipwreck of the Minotaur by J.M.W. Turner, 1805 • Femme nue assise by Pablo Picasso, 1910 • Composition VII by Wassily Kandinsky, 1913 • Der Schrei by Edvard Munch, 1893 • The Starry Night by Vincent van Gogh, 1889.

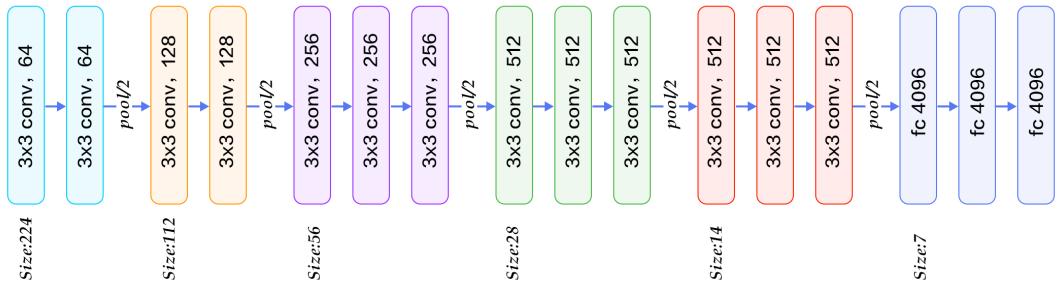


Figure 2.10: The structure of the VGG 19 network. It is for its most convolutional NN with 224×224 input size and with some downsampling layers which reduce the first two dimensions of the tensors along with the information flux of the network. At the very end of the architecture, there are three subsequent fully connected layers, responsible for the actual classification based on the features extracted from the previous layers.

1036 in equation (1.5). It can be computed between any couple of same-sized object,
 1037 hence also between the results of the same feature maps on the images X and C
 1038 at the same layer L :

$$L_{Cont} = \|F_{XL}F_{CL}\|^2. \quad (2.17)$$

1039 In order to evaluate this content loss, it is necessary to insert a custom transparent⁵
 1040 layer directly after the convolution layer(s) that are being used to compute the
 1041 content distance.

1042 Style Loss

1043 The concept of *style* loss function is the true novelty introduced by [18]. This loss
 1044 function is implemented similarly to the content loss module, as it will act as a
 1045 transparent layer in the network. The computation of the style loss requires in
 1046 advanced the evaluation of the Gram matrix G_{XL} at a certain layer L . A Gram
 1047 matrix is a result of multiplying a given matrix by its transposed matrix. In this
 1048 case, the matrix to multiplicate is a reshaped version of the feature maps F_{XL} :
 1049 \hat{F}_{XL} , a $K \times N$ matrix, where K is the number of feature maps at layer L and N
 1050 is the length of any vectorized feature map F_{XL}^k . Furthermore, the Gram matrix
 1051 must be normalized by dividing each element by the total number of elements in the
 1052 matrix. The style distance is now computed using the mean square error between
 1053 G_{XL} and G_{SL} :

$$L_{Style} = \|G_{XL}G_{SL}\|^2. \quad (2.18)$$

1054 After the appropriate insertion of the loss-function evaluator layers, one last piece for
 1055 finalizing the model is the right choice of the gradient descent optimizer. As in [18] and
 1056 according to the Deep Learning community the optimizer which suite best this role is the
 1057 Limited Memory-BFGS [8],[39]. L-BFGS is an iterative algorithm in the family of quasi-
 1058 Newton methods that approximates the Broyden-Fletcher-Goldfarb-Shanno algorithm
 1059 (BFGS) using a limited amount of computer memory, and it is a popular choice when
 1060 estimating parameters of a non-linear differentiable scalar function.

1061 The final phase of the training process thus makes run the optimizer for hundreds of
 1062 epochs on X , C , and S , and reduce the two loss functions values acting on the network's
 1063 parameters. After the fine-tuning of the weights, the hybrid image is produced, as in
 1064 Figure 2.9.

1065 2.8 Working Environment

1066 In this section I will briefly describe the machine I used to develop my project and the
 1067 working environment I built. All the work has been done on my personal computer,

⁵A transparent layer is a layer that performs some operations, like evaluating a function on its input, but returns as output an unchanged copy of its input.

1068 mounting a **GNU/linux** operating system, in particular 18.04 LTS **Ubuntu** version. The
1069 computer mounts an Intel i7 core, 8 Gb of RAM beside an 2Gb NVidia 940MX GPU.
1070 All the Python libraries have been installed and harmonized in a virtual environment
1071 mounting **Python 3.7.6**. All the code produced during the development, the images,
1072 and the data produced have been collected in a devoted repository on GitHub [14], which
1073 is freely available.

1074 As a conclusion for this chapter I will re-collect all the references to the different
1075 Python libraries I used during the development of this work:

1076 **NumPy:** NumPy is the pillar of every scientific computation-oriented library. Is the most
1077 spread library for heavy multidimensional numerical computation, and it offers a
1078 broad variety of tools like random number generators, and pre-implemented linear
1079 algebra utilities [30].

1080 **SciPy:** The SciPy library is one of the core packages that make up the SciPy stack.
1081 It provides many user-friendly and efficient numerical routines, such as routines
1082 for numerical integration, interpolation, optimization, linear algebra, and statistics
1083 [45]. Two modules in particular from this library have covered an essential
1084 role in this project: the `SciPy.spatial.Voronoi` module for the computation
1085 of the 3D Voronoi decomposition, as mentioned in section 2.3, and the
1086 `SciPy.spatial.ConvexHull` module for the computation of 3D and 2D convex
1087 hulls (section 2.5).

1088 **PyTorch:** PyTorch is a rich ecosystem of tools and libraries geared toward Machine
1089 Learning and Deep Learning. The application of the style-transfer NN described
1090 in section 2.7 has required the use of this framework [31].

1091 **SALib:** The SALib is a library which collects many tools for the Sensitivity Analysis of
1092 parameters. In particular the `SALib.saltelli` submodel was used for the quasi-
1093 random numerical sampling in a three-dimensional box, and it has been described
1094 in section 2.4.

1095 **pnoise2:** pnoise2 contains many tools for the production of specific types of noise.
1096 The module `noise` was tweaked for the production of two-dimensional Perlin noise
1097 surfaces, and it has been introduced in section 2.6.

1098 **pyquaternion:** The `pyquaternion` library provides a framework for handling quaternions.
1099 It has been widely used in section 2.1 for the design of three-dimensional
1100 ramifications for handling multiple spatial rotations.

₁₁₀₁ **Chapter 3**

₁₁₀₂ **Tissue Models Development**

₁₁₀₃ The main goal of the present work, as stated before, is to recreate a three-dimensional
₁₁₀₄ virtual model of histological tissue as faithfully as possible and then, to perform planar
₁₁₀₅ sectioning on it to emulate virtually the traditional histological specimen preparation
₁₁₀₆ procedure. The creation of a model of such complex structures is definitely a high-
₁₁₀₇ level problem, and it has required a careful designing, made of subsequent stages of
₁₁₀₈ improvements. In section 3.1, I will describe all the necessary steps to create the model
₁₁₀₉ of a small region of pancreatic tissue, while in section 3.2 I will expose the steps I
₁₁₁₀ followed to build a model of dermal tissue. In section 3.3, instead, I will show the
₁₁₁₁ resulting synthetic images from the sectioning process performed on both the models
₁₁₁₂ and all the enrichments and processing necessary to give them the most realistic look I
₁₁₁₃ was able to recreate.

₁₁₁₄ **3.1 Pancreatic Tissue Model**

₁₁₁₅ In this first attempt of modelization from scratch the main focus was put on reflecting
₁₁₁₆ only the main structural features on the virtual specimens. Given the pancreatic tissue's
₁₁₁₇ organization, described back in section 1.1.1, the first features I decided to put emphasis
₁₁₁₈ on were: 1) The iterative (with a fractal-like behavior) ramification of blood vessels for
₁₁₁₉ the irrigation of glandular acinus, 2) The space-filling distribution of acinus in the tissue,
₁₁₂₀ in fact, we expect a homogeneous density in the organ and to not see *holes* at all inside
₁₁₂₁ it. In this section I will describe step by step all the process I followed to create the
₁₁₂₂ model of a portion of pancreatic tissue, and all the interesting pitfalls I overcame.

₁₁₂₃ **3.1.1 2D Ramification**

₁₁₂₄ The first step was took in two dimensions, and it was the choice of the right *structure*
₁₁₂₅ to emulate the ramification of blood vessels in pancreatic tissue. The choice fell on a

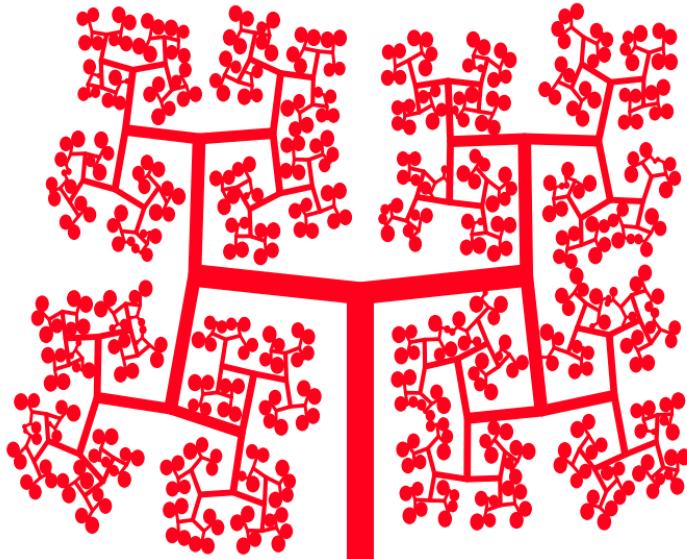


Figure 3.1: The development of the simple ramification in Figure 2.1, with some features to give it a more realistic look, like progressive thickness, angular noise in bifurcation and spheres at free ends of the ramification. The image is made using the tools exposed in section 2.2.

particular parametric L-system, as the one shown in Figure 2.1, in section 2.2. This structure is made of an iterative bifurcation of gradually shorter segments, with an angle of $\pm 85^\circ$ respect the main direction. For a start I added some features to give a more realistic look to the structure, which are all well represented in Figure 3.1:

- A progressive thickness of the bifurcation's segments, starting from a thick main branch that dwindle every junction. The idea is that the main blood vessel becomes gradually smaller becoming capillaries for single-cell irrigation.
- A progressive randomness in the angular deflection at every fork. Perfectly repeated angles are almost nonexistent in nature, so I decided to introduce an increasing indetermination in the angle of bifurcation from the main branch to the free ends of the structures' branches.
- Spheres at the ends of each branch, which acts as glandular acini. The maximum radius is comparable to the length of the final segments.
- A mechanism to avoid self-superimposition between branches and spheres. After the insertion of noise, the cumulative effect on the final segments might lead different branches to intersect. This is clearly a paradoxical situation, as real tissues while growing naturally occupy the space in a gradual way.

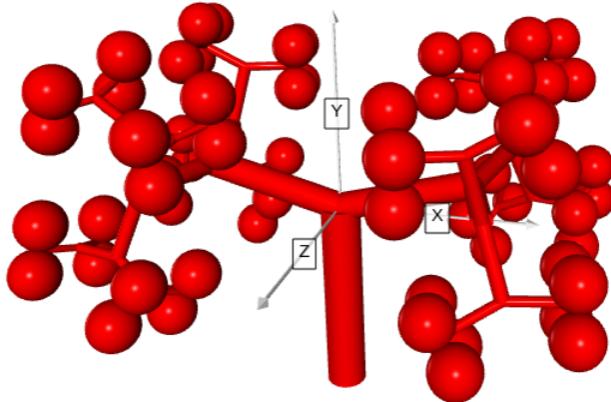


Figure 3.2: The three-dimensional expansion of the 2D ramification in Figure 3.1.

1143 To produce the specific image in Figure 3.1 I used a particular setting of the tool
 1144 described in section 2.2, which have a greatly wider range of customization and could be
 1145 used to create many other different structures to the need.

1146 3.1.2 Expansion to 3D

1147 The successive step I followed was to expand this structure in three dimensions and fill
 1148 the space in each of the three directions. The idea to evolve the structure in Figure 3.1 is
 1149 simply to twist of 90° the ramification at every junction point, in such a way to exit the
 1150 previous belonging plane. However, putting into practice this development has not been
 1151 easy. The organization of the structure in a 3D space requires an appropriate system
 1152 of reference for handling subsequent rotations in three dimensions. The best option for
 1153 handling relative 3D rotations, often used in computer graphics and every kind of 3D
 1154 modelization, are quaternions, as shown in section 2.1.

1155 In this new structure, segments are replaced with cylinders, and circles are replaced
 1156 with spheres. At every bifurcation to every cylinder are applied the following transfor-
 1157 mations:

- 1158 • a contraction in its extensions, regulated by an adjustable parameter R .
- 1159 • the usual deviation of $\pm 85^\circ$ respect to the direction of the parent branch.
- 1160 • a 90° specific rotation along the axis of its parent branch.

1161 The result of this procedure is a 3D ramification like the one in Figure 3.2, in which
 1162 we can recognize a good coverage of the space defined by the structure's boundaries and

immediate relation with the 2D structure in Figure 3.1. It should be noted that, in the further refinements of the model from now on, there won't be present the progressive angular indetermination on the direction of branches. Although it is a feature already implemented and working, it requires efficient control to avoid reciprocal overlapping between elements to produce a realistic structure. This second element has not been already developed and it would certainly enrich the representative power of the model.

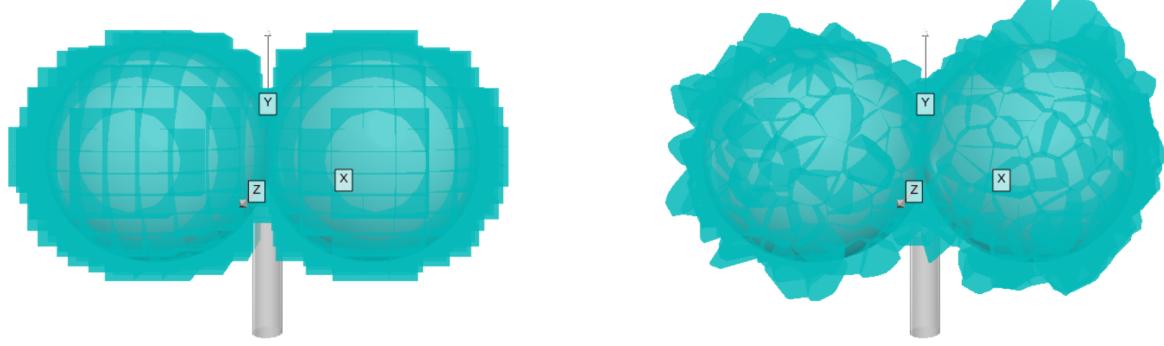
As for the 2D ramification the production of this structure has required the implementation of a tool for the 3D generation with a greatly wider power, able to produce almost any type of three-dimensional iterative structure after the right adjustment, and with a high degree of customization. It is necessary to mention the fundamental tool which allowed me to accomplish this step of the development, which is the `Python` library `VPython`: a library for 3D graphics visualization. This library allows a convenient and powerful interface to draw many types of objects and to move them around in space, which has been priceless to orient my self in three dimensions while developing the model and to produce all the 3D images visible in this work.

3.1.3 Subdivision in Cells

Once the 3D backbone of the pancreatic tissue blood vessels ramification system has taken shape, the next step was to embed all this structure in a spatial partitioning process, to create the subdivision into single cells. To perform this important task I used a 3D Voronoi decomposition, as shown in section 2.3. Depending on the choice of the starting points, the Voronoi tessellation could be an excellent item to recreate individual cells because it could guarantee some important properties: all the regions are convex, adjacent, with similar size and volume, with different shapes, and without holes. These have been chosen as the most significant properties to be reflected in the first modelization of cells.

As shown in section 2.3, the Voronoi decomposition strongly depends on the choice of the starting point. Points spread uniformly on a 3D regular lattice will produce a series of parallelepipeds repeated in the space. An example of uniform tessellation is shown in Figure 3.3a. On the other hand, a decomposition based on a quasi-random generated point can present all the good properties we mentioned before, including the diversity in shapes. In Figure 3.3b is shown an example of a Voronoi decomposition based on points sampled in a 3D with the Saltelli algorithm, in reference to section 2.4. Regardless of the points sampling technique, the boundaries of the sampling 3D box have been chosen to loosely contain the ramification.

There are tough some delicate considerations to be highlighted about the decomposition procedure. The first regards the most external pieces of the decomposition. Whilst the internal pieces are neatly bounded and defined, the most external layer instead is made on unbounded regions, which extend themselves to infinity. Those regions have clearly to be rejected, as it would be absurd for a cell to have an infinity volume. Typi-



(a) Regular lattice.

(b) Sampling with Saltelli algorithm.

Figure 3.3: Comparison between two Voronoi decompositions. The first (left) is created from a regular lattice of starting points, and every piece is exactly equal to all the others, creating a regular subdivision of the space. The second (right) is created instead from a sampling made following the Saltelli quasi-random algorithm. The pieces are all different in shape, but they all have similar sizes and volumes. In this pictures in particular have been shown only the pieces of the tessellation which lie in correspondence to the boundaries of the spheres underneath. While watching this picture one should immagine the decomposition extended similarly in all the space around the ramification, within certain boundaries, which loosely contains the structure. This limitation was necessary to enhance the interpretability of the image.

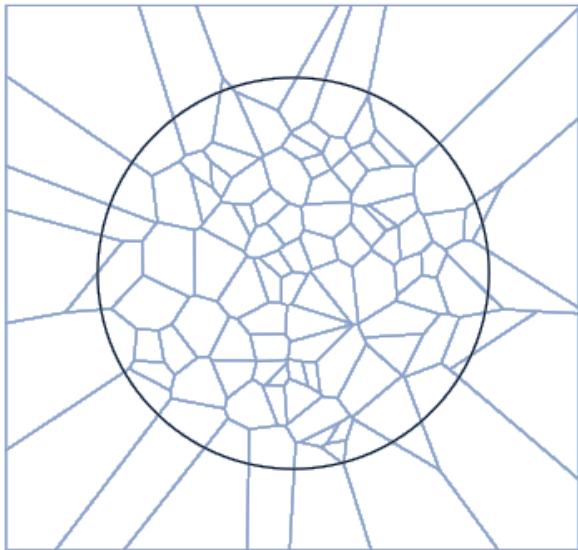


Figure 3.4: Example of circular cropping in a 2D Voronoi decomposition: all the regions which intersect the circumference have to be resized.

1202 cally those unbounded regions are resized in order to adhere to some limiting boundaries,
 1203 with an operation known as *cropping*. In Figure 3.4 is shown an example of circular crop-
 1204 ping in a 2D Voronoi decomposition: all the regions which intersect the circumference
 1205 have to be resized.

1206 The cropping operation in 3D is extremely complex, tough. Thus, a more simple and
 1207 efficient, yet less elegant, technique has been used. Instead of resizing the regions which
 1208 lie on the boundaries of the sampling region, those regions have directly been rejected.
 1209 This process is really fast and it does not lead to any danger of representativity loss if
 1210 the boundaries are loose enough and if the density of sampling is not too low.

1211 The other important consideration regards the density of sampling points. Increasing
 1212 the number of points to be extracted from the same volume automatically the number
 1213 of cells in the box will rise, and in contrast, their relative dimension will decrease. This
 1214 is a key element of the model: a too rarified decomposition would not bee able to reflect
 1215 the complexity of the structure underneath, but a too crowd decomposition on the other
 1216 side would lead to an unrealistic dimension of the cells in the tissue. Furthermore, this
 1217 parameter has a huge influence on the computing time necessary to generate the model
 1218 and to process it for the sectioning as will be shown in section 3.3. In almost all the
 1219 applications so far, the density parameter has been tuned by eye, with a trial and error
 1220 procedure. Although, a more rigorous way to adjust this parameter would be to consider
 1221 the average dimension of the cells and make some microanatomical considerations to

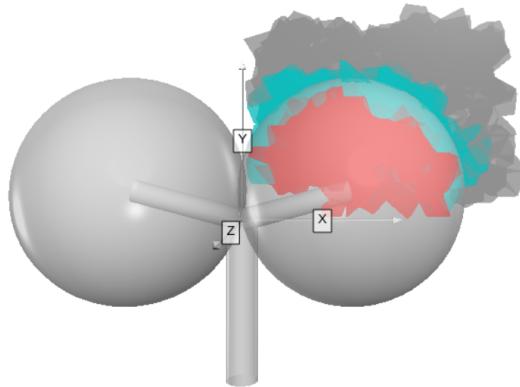


Figure 3.5: Portion of the complete Voronoi decomposition, showing the three different classes of cell in three different colors: the internal cells in red, those on the boundaries in turquoise and the external cells in gray.

1222 define the correct relative dimensions. The measure of the volume¹ of the decomposition's
 1223 regions is an accessible parameter, thus an easy way to estimate the average linear
 1224 dimension of the cells can be to approximate all the cells to cuboid seeing their volumes
 1225 as $V \approx L^3$. Averaging all the L measures an estimate \hat{L} can be done. This average
 1226 length may be compared to the length of the blood vessel ramification, allowing a good
 1227 reference tool.

1228 3.1.4 Cells Identity Assignment

1229 The great power of creating all the models virtually is to know exactly the identity of
 1230 every point in the structure. Although, This identity has to be reflected at the cellular
 1231 level, assigning to every region a label. Imagining the Voronoi decomposition represented
 1232 in Figure 3.3b extended to the entire box containing the ramification, good discrimination
 1233 would distinguish three classes of cells: those which lie completely inside a sphere, those
 1234 which lie completely outside a sphere, and those which lie on the boundaries of a sphere.
 1235 In Figure 3.5 is shown a portion of the complete decomposition where the three classes of
 1236 cells are reported with different colors: the internal cells in red, those on the boundaries
 1237 in turquoise, and the external cells in gray.

1238 In this particular case to find the relative position between every sphere in the struc-
 1239 ture and each cell it has been used a test on the proximity between the spheres' centers
 1240 and the vertices of every polyhedral cell. If all the vertices of a region lie within a dis-
 1241 tance lower than the radius from the center of the same sphere then that region can be
 1242 said to be an internal one. If none of the vertices lie within the radius distance from

¹The volume is expressed in the same arbitrary length unit of measures used during the ramification structure. This allows a coherent reference tool.

any center then that region is said to be external. In any other case, the region is said to be on the boundaries of some sphere, and this third label is assigned to it. As could be imagined the number of cells inside the volume can grow very quickly, and in the more rich ramifications also the number of spheres could be high. If we think that any polyhedron has a number of vertices of the order of 20/30 then it is clear that the number of distance evaluations could grow very quickly, requiring some relevant computational power in the more extended simulations. In order to optimize this computation, I decided to use a python implementation of a K-dimensional Tree, which is a space-partitioning data structure especially suited for fast and optimized computation of distances [6]. A K-d Tree is an algorithm that iteratively binary splits the space: every node of the three could be thought as a splitting $(k - 1)$ -hyperplane dividing the space into two semi-hyperspace. The result is an optimized algorithm for repeated distance evaluations. As for many other tools, in my code I used a pre-implemented module `KDTree` from the `Scipy` library.

This procedure of labeling the regions is completely customizable, and it should be adapted to the specific application. By the way, the principle will always be to perform some sort of spatial consideration respect to the primary structure and assign all the interesting labels accordingly to the cells in the volume.

After labeling the cells in the decomposition the model is considered complete. Every enrichment to the structure should be reflected in some type of label for the cells, which are chosen as the fundamental unit in the model. As we will see in section 3.3 during the sectioning process in the produced image will be printed mainly the identity of the cells, hence any detail on a finer scale in the model would not be conveyed properly on the final image.

3.2 Dermal Tissue Model

The modeling of the dermal tissue has followed analogous schedule respect to the previous model, hence many procedures and considerations have been repeated. The main target for this model was to recreate the stratification of different specific tissues in the section of a dermal sample. As clearly visible in Figure 1.2 in a dermal histological specimen one can distinguish the lighter and wider region of proper *dermal* tissue, underneath a more shallow and darker region of *epidermal* tissue. It is very interesting the boundary between those two regions, which can be seen as an irregular and smooth surface, populated of dermal lobes. On the other side of the epidermal layer lies a layer of *keratin*, with a smooth and regular boundary surface. Keratin is a family of fibrous structural proteins known as scleroproteins, a key structural material for hair, nails, and the outer layer of skin. The upper white region in Figure 1.2 instead is the support for the samples to perform optical analysis, and has no histological meaning.

1) Stratified Structure

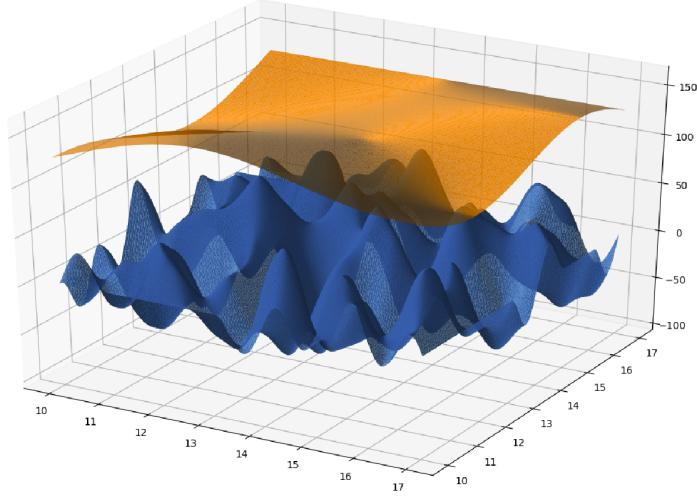


Figure 3.6: Pictures of two different Perlin noise surfaces used to separate dermal from epidermal tissue (blue) and epidermal from keratine layer (orange). The two surface are made by the same Perlin noise function, but the latter is stretched and compressed in order to have a more regular behavior.

In order to represent faithfully the stratification of different tissue layers, I decided to use one flat plane to separate epidermal tissue and the keratine layer and two boundary surfaces modulated by a Perlin noise function on different scales for the other two separations, as shown in Figure 3.6. As stated in section 2.6, after some easy customization the generation of different Perlin noise surfaces is easy and straightforward. By the way, to achieve the regularity and the smoothness of the orange surface in Figure 3.6 it was needed a horizontal stretching.

Following the scale of the image, the standard blue surface is created in a 7×7 square, while the orange surface has primarily been generated in a 1×1 square, and then has been stretched to cover the same 7×7 square, multiplying the values in its grid points respectively for the stretching factors $R_x = R_y = \frac{1}{7}$. In this primary structure, there are many important parameters defining the surfaces, like the distance between the two surfaces' average values and the amplitudes of the peaks and the valleys of the surfaces. In its standard version the Perlin noise covers the $[-1; 1]$ range, but with a simple multiplication for an amplitude factor A_S the values can be adjusted. Those particular values have been adjusted after some tries to recreate the proportions typical of a real specimen. To sum up, each one of the two surfaces is stored as a discretized three-dimensional array, or better as an array of 3-tuples in the form $(x, y, f_{(x,y)})$, one tuple for every (x, y) node of the grid, while the discretization grid was the same for both the surfaces.

2) Subdivision in Cells

1302 The subdivision in cells of the volume containing the structure has followed the
1303 exact same steps described in section 3.1, hence in this paragraph I will shortly
1304 resume the process. The first step is the definition of a suitable volume containing
1305 the structure. Then is the time for the generation of the decomposition's starting
1306 points according to a quasi-random number generation technique, as described in
1307 section 2.4. Afterward, the points are used as a base for the decomposition, and
1308 all the cells with undefined boundaries are rejected

1309 **3) Cells Identity Assignment**

1310 The identity assignment procedure instead has been customized for this particular
1311 application. In this model there are no *boundary* cells as the one lying on the
1312 spherical surfaces in the pancreatic model, thus there is no need for a test on the
1313 position of each vertex of every cell. In this model, the starting point for every
1314 Voronoi region has been used as a reference, and its relative position respect to the
1315 boundary surfaces was the discriminating factor for assessing the identity. In order
1316 to perform a coherent test on the relative position between the regions' center and
1317 the boundaries surfaces the positions of all the centers had to be discretized on
1318 the same grid onto which were defined the surfaces. The comparison with the flat
1319 horizontal plane defining the boundary between epidermal tissue and the keratine
1320 layer instead was simply a test on the z coordinate of the point: $z = f_{(x,y)} \leq \hat{z}_{plane}$.
1321 The result of this procedure is that every region is assigned a label corresponding
1322 to the belonging tissue layer: D for dermal, E for epidermal, K for keratine, and
1323 V , which stands for *void* for the white empty space above the sample.

1324 In this model, as in section 3.1, after the assignment of the identity to all the cells in
1325 the volume the modelization process is considered complete. Any sort of improvement
1326 and enrichment should be inserted during the structure designing phase and should be
1327 linked to an identity label.

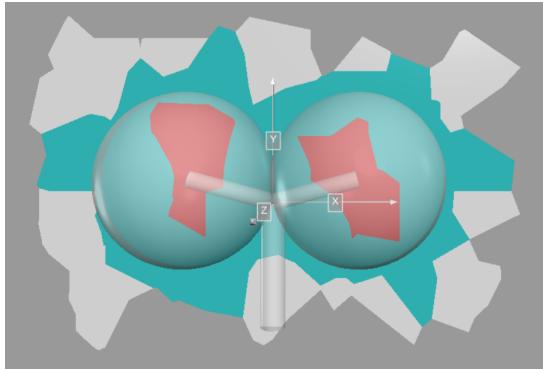
1328 3.3 Synthetic Images Production

1329 After the model is complete we have a three-dimensional representation of the tissue
1330 under study. The aim of the work is though to produce synthetic images from that
1331 structure, more precisely we want pairs of images composed of the synthetic-histological
1332 images and its related segmentation mask. The transition from 3D structure to a 2D
1333 image is the last step in the process, and it is inspired by the actual traditional technique
1334 for the preparation of the histological specimen, as described back in section 1.2.1. As
1335 the biopsy sample is treated and then sectioned with the microtome, the virtual model
1336 is sectioned in a random direction, producing an image representing the slice. This
1337 first image contains all the information of the section but its appearance is completely
1338 arbitrary and its look has nothing to share with a realistic sample. The original slice
1339 then acts perfectly as a segmentation mask, but some careful and dramatic makeover is
1340 needed to produce the final realistic looking image. In this section I will describe the
1341 general procedure to produce virtual slices from the two 3D virtual models described
1342 before in section 3 and the technique used to edit the images and give them the desired
1343 appearance.

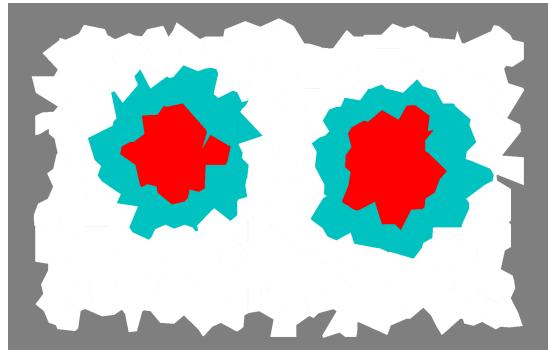
1344 3.3.1 Sectioning Process

1345 For any model created following the general procedure described in 3, even more so for
1346 the two particular models of pancreatic and dermal tissue, the sectioning process will
1347 be almost the same, and it will rely mainly on the algorithm for the general section of
1348 polyhedron described back in section 2.5. As stated before the models are essentially
1349 composed by labeled polyhedron spatially organized in a 3D volume. The ordered section
1350 of each polyhedron will yield all the polygons that shall be assembled in the final section.
1351 In Figure 3.7a is shown the three-dimensional representation of the section of a simple
1352 ramification, as the one in Figure 3.5. All the polygons that compose the section are
1353 drawn with the color correspondent to their label, following the same idea of Figure 3.5.
1354 In Figure 3.7b is printed the final result of the sectioning algorithm applied to the model,
1355 which will be the segmentation mask in the single pair of synthetic images. The colors
1356 in the produced slice match the colors used for the different identities in the 3D model.

1357 The simplest, yet over-abundant, way to proceed is to create the model in its entirety
1358 and subsequently choosing the sectioning plane. Afterward, it is necessary to select
1359 only the regions that intersect the plane and section them all. Actually, the test on the
1360 intersection passes through the check on the relative position of the polyhedron's vertices
1361 respect to the sectioning plane: if all the vertices lie on the same semi-space then the
1362 intersection would be null and the polyhedron is not of interest for that particular section.
1363 This procedure is exactly the first step of the algorithm in 2.5, thus the filtering on the
1364 regions is actually made during construction for optimization. The alternative method
1365 could be to choose in the first place the direction of the sectioning plane, and in second



(a) 2D polygonal sections represented in a 3D environment.



(b) The final section image, produced directly on a planar picture, skipping completely the 3D representation.

Figure 3.7: In this Figure is shown the idea of the correspondence between the section of a simple structure in the space and the correspondent section. The correspondence is not perfect for representation requirements, in fact the two images even if very similar are produced with two completely different methods. At the left an image of 2D polygonal section embeded in a 3D space, made using 3D visualization tools. At the right a simple image produced printing the polygons in a planar picture. Printing 2D polygons in a 3D space is much more complicated than one would think using the same tool used to produce the other images like Figure 3.5 and 3.3. This choice has been done for the sake of the overall homogeneity in pictures style.

1366 place to generate the model's decomposition only in the volume adjacent to that plane.
 1367 This method enables the sparing of a good amount of computation, without any negative
 1368 impact on the final result. The only delicate step is the choice of a wide enough region
 1369 of space around the plane, which doesn't compromise the representation.

1370 As a guarantee for richness and diversification among the images, there is the need
 1371 for some degree of controlled randomness in the sectioning process, for example in the
 1372 determination of the sectioning plane direction. All the sectioning process is then based
 1373 on a single starting *seed*, which determines the direction of the sectioning plane in a
 1374 deterministic way, and all the rest of the model is generated as a consequence. In this
 1375 way, all the possible angulations are equally probable and will be sampled in view of
 1376 multiple applications of this process. In Figure 3.8 are shown two different sections,
 1377 along two random planes on two simple ramified structures with four spheres.

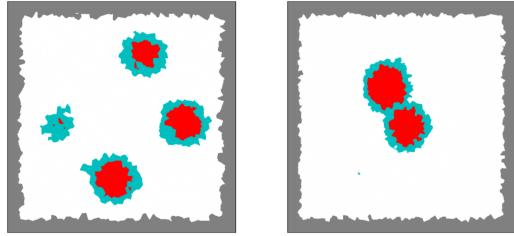


Figure 3.8: Two different section, along two random planes on two simple ramified structure with four spheres.

1378 **3.3.2 Appearance Makeover**

1379 After the application of the sectioning algorithm of the previous section, the image
 1380 which will act as a segmentation mask is ready. The last and most complex task that
 1381 remains is to transform the image and to give it a realistic look. I tried many different
 1382 transformations, more or less complicated, and there was not a final decision on which is
 1383 the best blend of them. In this section, I will describe them as an arsenal of possibilities
 1384 and show their impact on the images.

1385 **Color Palette**

1386 The first correction to do to the images will inevitably be a change in the colors
 1387 of the image. Gray, Turquoise, and Red are the perfect choice for label-colors but
 1388 act poorly as physiological colors. In Figure 3.9 is shown an example of an image
 1389 produced re-mapping the colors to a new palette, inspired to the coloring of the
 1390 real specimen in Figure 1.1 and 1.4, given by the traditional hematoxylin and eosin
 1391 staining process.

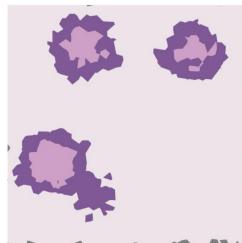


Figure 3.9: Example of images produced re-mapping the colors with a color palette inspired to a real H&E stained histological sample.

1392 **Nuclei Projection**

1393 Another fundamental processing needed was the projection of cells' nuclei on the
 1394 image. Usually, nuclei are clearly visible in histological samples and guide the

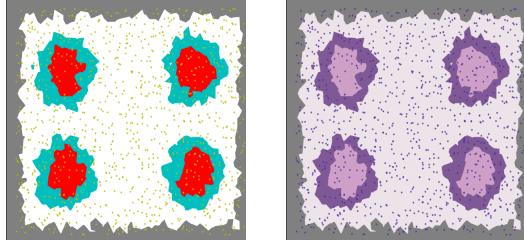


Figure 3.10: Nuclei projection on the image: (left) in yellow in the segmentation mask and (right) in purple in the image under makeover.

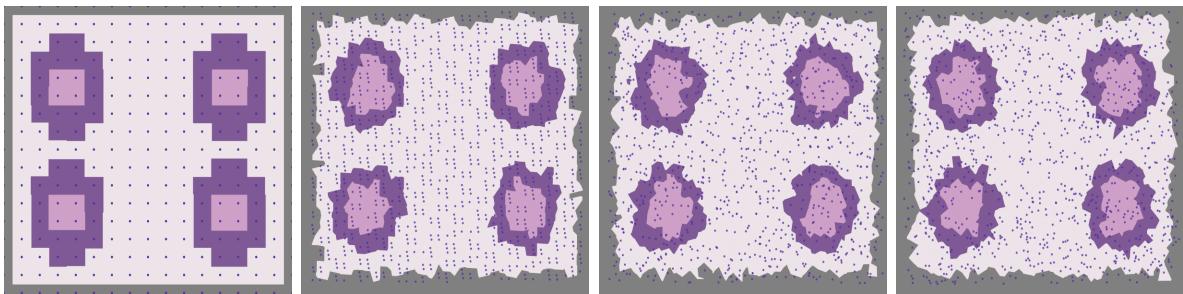


Figure 3.11: Four different sections produced with the same density on cells but with four different method for the sampling of starting decomposition's ponits (from left to right): • points sampled on a regular lattice, • sampling following a simple recursion sequence as the one in equation (2.12), • following the saltelli algorithm, • following a fully-random distribution.

analysis allowing to detect individual cells in the specimen. As a reference for the nucleus position the starting point of every polyhedral region has been used and projected on the sectioning plane as a little dark circle. The diameter of those circles as been chosen to be a submultiple (10%) of the linear estimated dimension \hat{L} of the cells in the decomposition².

Nuclei projection, among the other things, is an excellent tool to perceive the different effects obtained with different choices of quasi-random distributions or fully-random distributions (with reference to section 2.4). The different impact on the overall image is huge, and it really changes the overall sense of the image. In Figure 3.11 are reported four different sections, produced with the same density on cells but with four different methods for the sampling of starting decomposition's points.

1407 **Boundaries Projection**

²Following the same logic of step 3 of section 3.1.

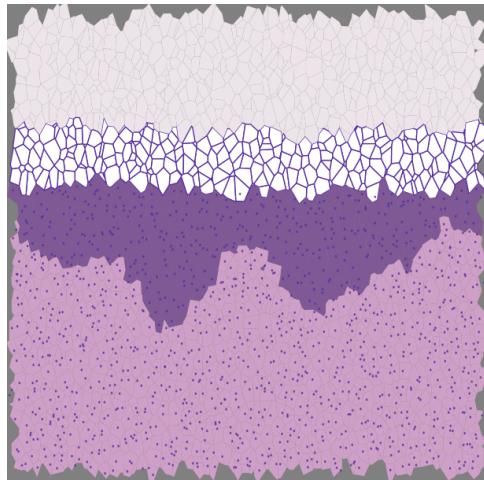


Figure 3.12: Example of images produced re-mapping the colors with a color palette inspired to a real H&E stained histological sample.

On the same wave of the previous tool, another operation that can help the appearance of an image is the projection of the boundaries of each (or just a part) of the polygonal sections. The drawing can be clearly tuned and customized depending on the specific necessities. In Figure 3.12 is shown an example of a section on the dermal tissue model in 3.2, in which the boundaries of all the cells have been lightly marked, and the boundaries of the cells in the keratine layer have been heavily marked instead.

Blurring Effects

In all the images produced so far the boundaries between polygonal sections are perfectly sharp and without any smudge. To give a more realistic feeling to those pictures I tried different forms of blurring. As a first try, I applied a Gaussian blurring filter, which is an extremely common blurring operation in computer vision, which consists of a simple discrete convolution with a 2D Gaussian kernel. The effect is a regular and diffuse blur all over the image, as in Figure 3.13a. The second blurring effect I implemented was based instead on the averaging of parallel and adjacent slices on the same model. This method is inspired by the real sectioning technique (section 1.2.1), in which every slice is not an infinitesimal layer of matter, but a finite sample, which suffers from mechanical dragging during the process. The idea is that the average between three or more slices equally spaced above and below the *main* slice should recreate a realistic blurring effect. An example of an image produced with this process is shown in Figure 3.13b.

Perlin RGB Noise

A further attempt to give visual texture to the image was done using again the

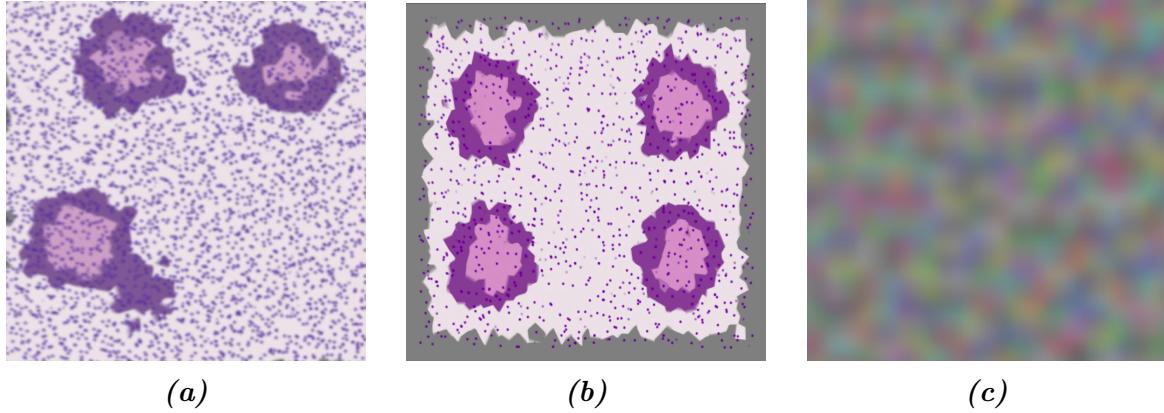


Figure 3.13: The two blurring effects used in this work: (left) A standard Gaussian-filter blur and (center) a specific blur introduced averaging adjacent parallel slices on the same image, (right) an example of RGB color noise built joining three different Perlin noise surfaces, one for each color channel.

Perlin noise, described in section 2.6. The idea is to create some fluctuation among the color channels of the image around the sharp values of the image produced by the sectioning algorithm. From a practical point of view, I created three different and independent Perlin noise surfaces, one for every color channel (Red, Green, and Blue), and added them to create an RGB noise on the image. An example of the resulting image is shown in Figure 3.13c.

Style Transfer

This last tool I will describe is the most sophisticated so far. It consists of the application of a style-transfer neural network (STNN) on the image obtained through the sectioning process, for the implantation of the visual texture from a real sample of the corresponding tissue. Style-transfer NNs, and their functioning, have been described in detail in section 2.7, and here I will cover just the particular applications on the two type of section produced.

The first manipulation I report is the one on a section from the pancreatic tissue model. The image of which to conserve the visual content is a section with some simple pre-processing picked from the ones described before (Figure 3.14a): a more accurate color palette, the projection of nuclei, and the average on five adjacent slices. The image from which to pick the style, thus the visual texture, is a portion of an actual histological sample of the pancreas, and it is shown in Figure 3.14b. The application of the STNN yields a hybrid image, shown in Figure 3.14c. The second application was made on a section on dermal tissue. The three content, style and styled images are shown respectively in Figure 3.15a, 3.15b, and 3.15c.

In Figures 3.14, and 3.15 are reported the best results among all the different tests

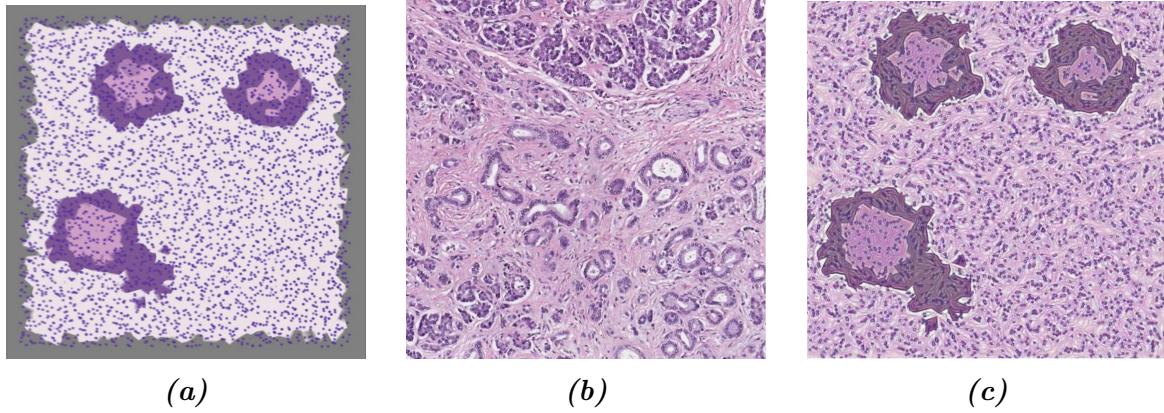


Figure 3.14: Application of the style-transfer NN on a section of the pancreatic tissue model: 3.14a the content image, 3.14b the style image, 3.14c the hybrid resulting image.

made on the sections. I made different tries on the same image with different processing before the manipulation with the STNN, to see the impact of the different adjustment on the resulting styled image. It turned out that the presence of nuclei is essential to give a homogeneous texture to the image and avoid unrealistic artifacts. On the other hand, the choice of the color palette has a way lighter effect than what one would think: the model yields almost the same result with a grey-levels image or with any other palette.

It is interesting to notice the timing cost of this style transfer operation. While all the other manipulation described in this chapter requires a very short time (seconds) to be applied, and are in practice *instantaneous*, the transfer of style is a way more robust operation, which implies the finalization of the training of a pre-trained neural network. On a computer with the technical specification described in section 2.8 this operation instead took minutes, which is a time two full orders of magnitude greater.

It should be noted that the presented results are obtained from the application of a pre-trained STNN model. The development of a specialized model for histological texture transfer could improve extremely the ability to produce realistic images, way further the present results.

One single complete application of the process consists then in the generation of a tissue model, in the sectioning along a random section plane, and in the processing of the image, in order to produce the pair of ground-truth image and the synthetic histological image. The target is to apply over and over this process to collect the necessary amount of images and constitute an entire dataset. An important feature to have for the process is thus a complete automatization, in order to scale up the

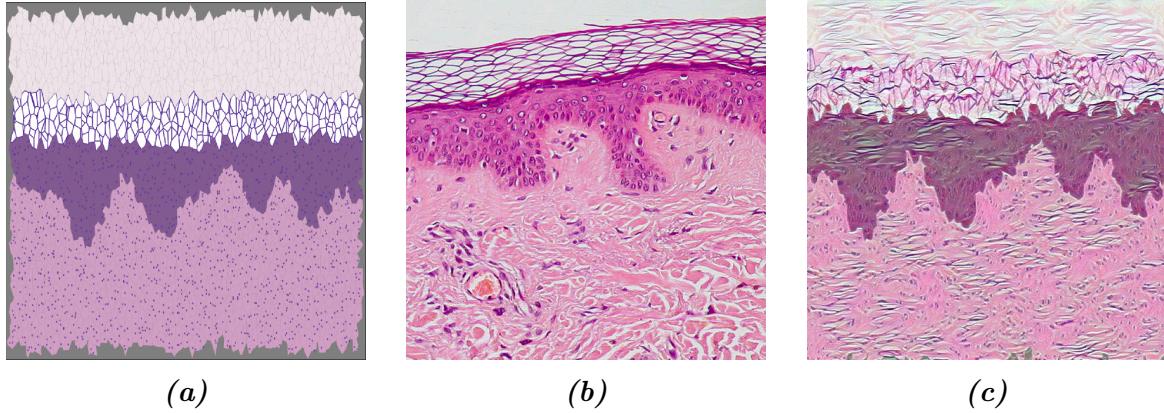


Figure 3.15: Application of the style-transfer NN on a section of the dermal tissue model: 3.15a the content image, 3.15b the style image, 3.15c the hybrid resulting image.

1478 generation of images, possibly even in parallel computation. For this purpose, I created a
 1479 pipeline workflow interface for the image generation, with an automatized harmonization
 1480 of every piece of the process. The generation now requires just to fill a configuration
 1481 file in which writes all the specific characteristics of the images: the type of structure,
 1482 its features, the desired processing on the images, and eventually the random seeds for
 1483 a supervised generation. In Figure 3.16 is reported a small scale example of a dataset
 1484 produced with multiple automatized applications of the generation tool on a ramified
 1485 structure inspired to a pancreatic tissue model. It is clear the correspondence between
 1486 segmentation mask and synthetic histological images, and the diversification given by
 1487 the supervised randomness on the generation.

1488 The actual tool used for the set up of a working pipeline was the **Snakemake** [24]
 1489 workflow management system, which is a python-based tool to create reproducible and
 1490 scalable data analyses.

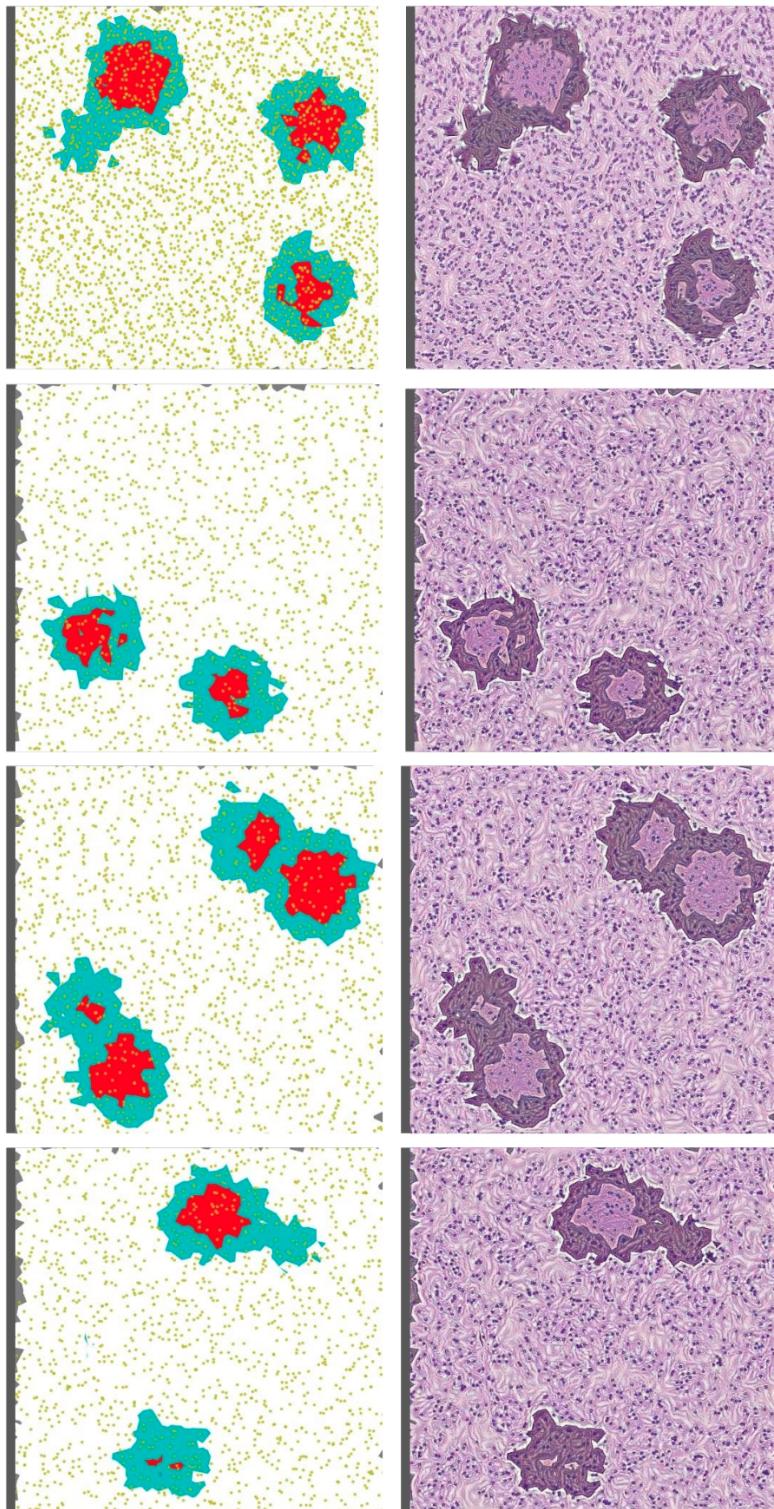


Figure 3.16: Small scale example of dataset produced with multiple automated applications of the generation tool on a ramified structure inspired to a pancreatic tissue model.

¹⁴⁹¹ Conclusions

¹⁴⁹² In this project, I face the problem of synthetic histological image generation for the
¹⁴⁹³ purpose of training Neural Networks for the segmentation of real histological images. The
¹⁴⁹⁴ manual analysis of histological specimen is a complex, time consuming, and expensive
¹⁴⁹⁵ task and nevertheless it is a pillar of countless diagnostic techniques. Any form of
¹⁴⁹⁶ support for this procedure hence is welcome and endorsed by the health-care system. In
¹⁴⁹⁷ particular, in this work, I focus on the problem of histological specimens segmentation.
¹⁴⁹⁸ The most advanced algorithms for image segmentation are based on Deep Learning
¹⁴⁹⁹ and requires the training of extensive and complex neural networks. One of the toughest
¹⁵⁰⁰ hurdles to overcome for the training of those NN is the abundance and the quality of pre-
¹⁵⁰¹ labeled examples of segmentation on real histological samples. The collection of hundreds
¹⁵⁰² of hand-labeled histological samples, with pixel-level precision, is virtually impossible.
¹⁵⁰³ This work thus proposes a methodology to generate, in a completely automatic way,
¹⁵⁰⁴ synthetic pre-labeled histological-like images, that can be used as training material for
¹⁵⁰⁵ a NN.

¹⁵⁰⁶ The method I propose consists of the recreation of the traditional histological speci-
¹⁵⁰⁷ mens' preparation, and it is based on the sectioning of a 3D virtual model of a region of
¹⁵⁰⁸ histological tissue. The virtual 3D model of a region of a particular type of human tissue
¹⁵⁰⁹ is built after physical and physiological considerations. The model is then subject to a
¹⁵¹⁰ virtual sectioning operation, which yields the synthetic sampling of the virtual tissue in
¹⁵¹¹ which the histological identity of every pixel is perfectly known. This first image will act
¹⁵¹² as a segmentation mask for a second, realistic image. In fact, on top of this first image
¹⁵¹³ are applied several aesthetical processing and refinements and the final product is the
¹⁵¹⁴ synthetic histological-like image. The pair made of the two images is perfectly suitable
¹⁵¹⁵ for the supervised learning of a NN oriented toward the segmentation of histological
¹⁵¹⁶ images. The production of each pair of images is completely automatic and it does not
¹⁵¹⁷ require the intervention of any human operator, it is thus a scalable process that can
¹⁵¹⁸ produce a great abundance of images. The quality of the images is directly connected
¹⁵¹⁹ to the richness and the quality of the model. The perfect modelization of a region of
¹⁵²⁰ tissue, let's say human pancreatic tissue, is by far out of reach for this, hence the rich-
¹⁵²¹ ness and the fidelity of the produced images are inevitably lower than the real sample.
¹⁵²² Nevertheless, the quality of the produced images is sufficient to perform the preliminary

1523 phase of the training of a NN following a training strategy known as curriculum learning.
1524 This learning process consists of giving the NN a copious quantity of lower complexity
1525 level example in the first instance, reserving the few and sophisticated real hand-labeled
1526 histological samples for the finalization of the training.

1527 The first chapter of this thesis is devoted to the contextualization of the present
1528 work. It is offered a description of how the histological samples are obtained after a
1529 tissue biopsy and how the digitalization process of the images works.

1530 The second chapter collects all the details of every less common technical tool I used
1531 during the design of this project. A brief theoretical introduction is proposed for every
1532 item besides the thorough description of its practical use. In this chapter, a section is
1533 devoted to the description of a general methodology for computing the 2D section of an
1534 arbitrary three-dimensional polyhedron. The algorithm here described has been devised
1535 and implemented all by my self, and then inserted in the workflow of the project. This
1536 is one of the key pieces for the automatization of the virtual tomography process, and it
1537 allows to connect the three dimensional model to the two-dimensional representation of
1538 a sampled section.

1539 The third chapter is the center of this work, and it contains the description of all
1540 the design choices, and the steps I followed for the development of the two human tissue
1541 models I propose: the first of pancreatic tissue and the second of dermal tissue. The first
1542 and second sections are dedicated to the description of the two proposed 3D virtual tissue
1543 models, which consists of different steps. The third section instead contains a thorough
1544 description of the method to perform the sectioning onto a virtual model and how to
1545 process the resulting images. The development has required the harmonization of many
1546 different technical aspects and mathematical tools and it results in a general methodology
1547 for the generation of synthetic histological images. The process passes first through the
1548 building of the target tissue's structure in a virtual environment. This structure is then
1549 embedded in a three-dimensional space decomposition which subdivide the volume in
1550 individual cells. Those cells are labeled in correspondence to their role in the model, and
1551 their identity is then perfectly encaptured by the virtual tomography procedure. The
1552 sectioning process is responsables for the production of synthetic images, which are then
1553 conveyed toward an aesthetical enrichment pipeline specialized for the particular target
1554 tissue. The product of any application of this process is a pair made of segmentation mask
1555 and the correpsonding synthetic histological-like image. This completely automatized
1556 procedure allow to built arbitrary large dataset for the training of NN, without the
1557 intervention of any human operator.

1558 The method for the generation of datasets of synthetic images I propose in my thesis
1559 work is a self-supporting project and it is formally consistent. By the way, there are
1560 many possibilities for improvement and enrichment for the project. One first aspect to
1561 strengthen could be the richness of the models: adding more elements in the structure,
1562 and refining their representation of the tissue at the cellular level. This would lead
1563 to a better quality of the synthetic images, that would assist the training of NN in

1564 more and different applications. Another aspect that lends itself to improvements in the
1565 development of a dedicated style transfer NN targeting the histological texture transfer,
1566 which could lead to interesting signs of progress in the quality of image generation. There
1567 is also the intention to perform an actual attempt of NN training on the images produced
1568 with this process. This would complete conceptually the idea underneath the project
1569 and would be an excellent opportunity to detect weaknesses and to draw up possible
1570 lines of development. The repeated application of the generation method would allow
1571 the building of entire datasets suitable for the training of DL-based models.

₁₅₇₂

Bibliography

- ₁₅₇₃ [1] Abien Fred Agarap. Deep learning using rectified linear units (relu), 2018.
- ₁₅₇₄ [2] Salah Alheejawi, Mrinal Mandal, Hongming Xu, Cheng Lu, Richard Berendt, and
₁₅₇₅ Naresh Jha. Deep learning-based histopathological image analysis for automated
₁₅₇₆ detection and staging of melanoma. In *Deep Learning Techniques for Biomedical*
₁₅₇₇ and *Health Informatics*, pages 237–265. Elsevier, 2020.
- ₁₅₇₈ [3] J. Alsayednoor and P. Harrison. Evaluating the performance of microstructure
₁₅₇₉ generation algorithms for 2-d foam-like representative volume elements. *Mechanics*
₁₅₈₀ of *Materials*, 98:44 – 58, 2016.
- ₁₅₈₁ [4] Hani A Alturkistani, Faris M Tashkandi, and Zuhair M Mohammedsaleh. Histolog-
₁₅₈₂ ical stains: A literature review and case study. *Global Journal of Health Science*,
₁₅₈₃ 8(3):72, June 2015.
- ₁₅₈₄ [5] Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curricu-
₁₅₈₅ lum learning. In *Proceedings of the 26th Annual International Conference on Ma-*
₁₅₈₆ *chine Learning*, ICML '09, page 4148, New York, NY, USA, 2009. Association for
₁₅₈₇ Computing Machinery.
- ₁₅₈₈ [6] Jon Louis Bentley. Multidimensional binary search trees used for associative search-
₁₅₈₉ ing. *Commun. ACM*, 18(9):509517, September 1975.
- ₁₅₉₀ [7] R. W. Brockett. Robotic manipulators and the product of exponentials formula.
₁₅₉₁ In P. A. Fuhrmann, editor, *Mathematical Theory of Networks and Systems*, pages
₁₅₉₂ 120–129, Berlin, Heidelberg, 1984. Springer Berlin Heidelberg.
- ₁₅₉₃ [8] C. G. BROYDEN. The Convergence of a Class of Double-rank Minimization Algo-
₁₅₉₄ rithms 1. General Considerations. *IMA Journal of Applied Mathematics*, 6(1):76–90,
₁₅₉₅ 03 1970.
- ₁₅₉₆ [9] Bassem Ben Cheikh, Catherine Bor-Angelier, and Daniel Racoceanu. A model of
₁₅₉₇ tumor architecture and spatial interactions with tumor microenvironment in breast
₁₅₉₈ carcinoma. In Metin N. Gurcan and John E. Tomaszewski, editors, *Medical Imaging*

- 1599 2017: *Digital Pathology*, volume 10140, pages 73 – 80. International Society for
 1600 Optics and Photonics, SPIE, 2017.
- 1601 [10] Anna Choromanska, Benjamin Cowen, Sadhana Kumaravel, Ronny Luss, Mattia
 1602 Rigotti, Irina Rish, Brian Kingsbury, Paolo DiAchille, Viatcheslav Gurev, Ravi
 1603 Tejwani, and Djallel Bouneffouf. Beyond backprop: Online alternating minimization
 1604 with auxiliary variables, 2018.
- 1605 [11] Siddharth Singh Chouhan, Ajay Kaul, and Uday Pratap Singh. Image segmentation
 1606 using computational intelligence techniques: Review. *Archives of Computational
 1607 Methods in Engineering*, 26(3):533–596, February 2018.
- 1608 [12] Marius Cordts, Mohamed Omran, Sebastian Ramos, Timo Rehfeld, Markus En-
 1609 zweiler, Rodrigo Benenson, Uwe Franke, Stefan Roth, and Bernt Schiele. The
 1610 cityscapes dataset for semantic urban scene understanding. *CoRR*, abs/1604.01685,
 1611 2016.
- 1612 [13] Angel Cruz-Roa, Ajay Basavanhally, Fabio Gonzlez, Hannah Gilmore, Michael Feld-
 1613 man, Shridar Ganesan, Natalie Shih, John Tomaszewski, and Anant Madabhushi.
 1614 Automatic detection of invasive ductal carcinoma in whole slide images with convo-
 1615 lutional neural networks. *Progress in Biomedical Optics and Imaging - Proceedings
 1616 of SPIE*, 9041, 02 2014.
- 1617 [14] Alessandro d'Agostino. Dataset generation for the training of neural networks ori-
 1618 ented toward histological image segmentation, april 2020.
- 1619 [15] J. Deng, W. Dong, R. Socher, L. Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale
 1620 hierarchical image database. In *2009 IEEE Conference on Computer Vision and
 1621 Pattern Recognition*, pages 248–255, 2009.
- 1622 [16] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei. ImageNet: A Large-
 1623 Scale Hierarchical Image Database. In *CVPR09*, 2009.
- 1624 [17] Mark Everingham, Luc Van Gool, Christopher KI Williams, John Winn, and An-
 1625 drew Zisserman. The pascal visual object classes (voc) challenge. *International
 1626 journal of computer vision*, 88(2):303–338, 2010.
- 1627 [18] Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge. A neural algorithm of
 1628 artistic style, 2015.
- 1629 [19] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural com-
 1630 putation*, 9(8):1735–1780, 1997.
- 1631 [20] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros. Image-to-image
 1632 translation with conditional adversarial networks, 2016.

- 1633 [21] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization,
1634 2014.
- 1635 [22] Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical
1636 report, 2009.
- 1637 [23] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with
1638 deep convolutional neural networks. In *Advances in neural information processing*
1639 *systems*, pages 1097–1105, 2012.
- 1640 [24] Johannes Kster and Sven Rahmann. Snakemakea scalable bioinformatics workflow
1641 engine. *Bioinformatics*, 28(19):2520–2522, 08 2012.
- 1642 [25] Aristid Lindenmayer. Mathematical models for cellular interactions in development
1643 ii. simple and branching filaments with two-sided inputs. *Journal of theoretical*
1644 *biology*, 18(3):300–315, 1968.
- 1645 [26] Daniel. Longnecker. Anatomy and histology of the pancreas, 2014.
- 1646 [27] Shervin Minaee, Yuri Boykov, Fatih Porikli, Antonio Plaza, Nasser Kehtarnavaz,
1647 and Demetri Terzopoulos. Image segmentation using deep learning: A survey, 2020.
- 1648 [28] Muhammad Niazi, Thomas Tavolara, Vidya Arole, Douglas Hartman, Liron Pan-
1649 tanowitz, and Metin Gurcan. Identifying tumor in pancreatic neuroendocrine neo-
1650 plasms from ki67 images using transfer learning. *PLOS ONE*, 13:e0195621, 04 2018.
- 1651 [29] Arild Nkland and Lars Hiller Eidnes. Training neural networks with local error
1652 signals, 2019.
- 1653 [30] Travis E Oliphant. *A guide to NumPy*, volume 1. Trelgol Publishing USA, 2006.
- 1654 [31] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gre-
1655 gory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban
1656 Desmaison, Andreas Kopf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan
1657 Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith
1658 Chintala. Pytorch: An imperative style, high-performance deep learning library,
1659 2019.
- 1660 [32] Allan Pinkus. Approximation theory of the mlp model in neural networks. *Acta*
1661 *Numerica*, 8:143195, 1999.
- 1662 [33] Prabu Ravindran, Adriana Costa, Richard Soares, and Alex C Wiedenhoeft. Clas-
1663 sification of cites-listed and other neotropical meliaceae wood images using convo-
1664 lutional neural networks. *Plant methods*, 14(1):1–10, 2018.

- 1665 [34] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional net-
1666 works for biomedical image segmentation. *CoRR*, abs/1505.04597, 2015.
- 1667 [35] Juan Rosai. Why microscopy will remain a cornerstone of surgical pathology. *Lab-*
1668 *oratory Investigation*, 87(5):403–408, April 2007.
- 1669 [36] Andrea Saltelli. Making best use of model evaluations to compute sensitivity indices.
1670 *Computer Physics Communications*, 145(2):280 – 297, 2002.
- 1671 [37] Andrea Saltelli, Paola Annoni, Ivano Azzini, Francesca Campolongo, Marco Ratto,
1672 and Stefano Tarantola. Variance based sensitivity analysis of model output. design
1673 and estimator for the total sensitivity index. *Computer Physics Communications*,
1674 181(2):259 – 270, 2010.
- 1675 [38] Caglar Senaras, Muhammad Khalid Khan Niazi, Berkman Sahiner, Michael P. Pen-
1676 nell, Gary Tozbikian, Gerard Lozanski, and Metin N. Gurcan. Optimized generation
1677 of high-resolution phantom images using cGAN: Application to quantification of ki67
1678 breast cancer images. *PLOS ONE*, 13(5):e0196846, May 2018.
- 1679 [39] David F Shanno. Conditioning of quasi-newton methods for function minimization.
1680 *Mathematics of computation*, 24(111):647–656, 1970.
- 1681 [40] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for
1682 large-scale image recognition, 2014.
- 1683 [41] Sandro Skansi. *Introduction to Deep Learning: From Logical Calculus to Artificial*
1684 *Intelligence*. Springer Publishing Company, Incorporated, 1st edition, 2018.
- 1685 [42] I.M. Sobol. Uniformly distributed sequences with an additional uniform prop-
1686 erty. *USSR Computational Mathematics and Mathematical Physics*, 16(5):236 –
1687 242, 1976.
- 1688 [43] I.M Sobol. Global sensitivity indices for nonlinear mathematical models and their
1689 monte carlo estimates. *Mathematics and Computers in Simulation*, 55(1):271 – 280,
1690 2001. The Second IMACS Seminar on Monte Carlo Methods.
- 1691 [44] M Titford. The long history of hematoxylin. *Biotechnic & Histochemistry*, 80(2):73–
1692 78, 2005.
- 1693 [45] Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, Matt Haberland, Tyler Reddy,
1694 David Cournapeau, Evgeni Burovski, Pearu Peterson, Warren Weckesser, Jonathan
1695 Bright, Stéfan J. van der Walt, Matthew Brett, Joshua Wilson, K. Jarrod Millman,
1696 Nikolay Mayorov, Andrew R. J. Nelson, Eric Jones, Robert Kern, Eric Larson,
1697 CJ Carey, İlhan Polat, Yu Feng, Eric W. Moore, Jake VanderPlas, Denis Laxalde,

1698 Josef Perktold, Robert Cimrman, Ian Henriksen, E. A. Quintero, Charles R Harris,
1699 Anne M. Archibald, Antônio H. Ribeiro, Fabian Pedregosa, Paul van Mulbregt,
1700 and SciPy 1. 0 Contributors. SciPy 1.0: Fundamental Algorithms for Scientific
1701 Computing in Python. *Nature Methods*, 17:261–272, 2020.

1702 [46] Georges Voronoi. Nouvelles applications des paramètres continus à la théorie des
1703 formes quadratiques. premier mémoire. sur quelques propriétés des formes quadra-
1704 tiques positives parfaites. *Journal für die reine und angewandte Mathematik (Crelles
1705 Journal)*, 1908:102 – 97.